UNIVERSITAT DE BARCELONA

MSc in Fundamental Principles of Data Science

Ethical Data Science

Bias and Discrimination

Jordi Vitrià

2020-2021

Definitions

In ML, ideally, one wants to choose a model that both accurately captures the **regularities in its training data**, but also **generalizes well** to unseen data.

Unfortunately, it is typically impossible to do both simultaneously. High-variance learning methods may be able to represent their training set well but are at risk of overfitting to noisy or unrepresentative training data. In contrast, algorithms with **high bias** typically produce simpler models that may fail to capture important regularities (i.e. underfit) in the data.

The **bias error** is an error from erroneous assumptions in the learning algorithm. High bias can cause an algorithm to miss the relevant relations between features and target outputs (underfitting).

Source: Wikipedia

Definitions

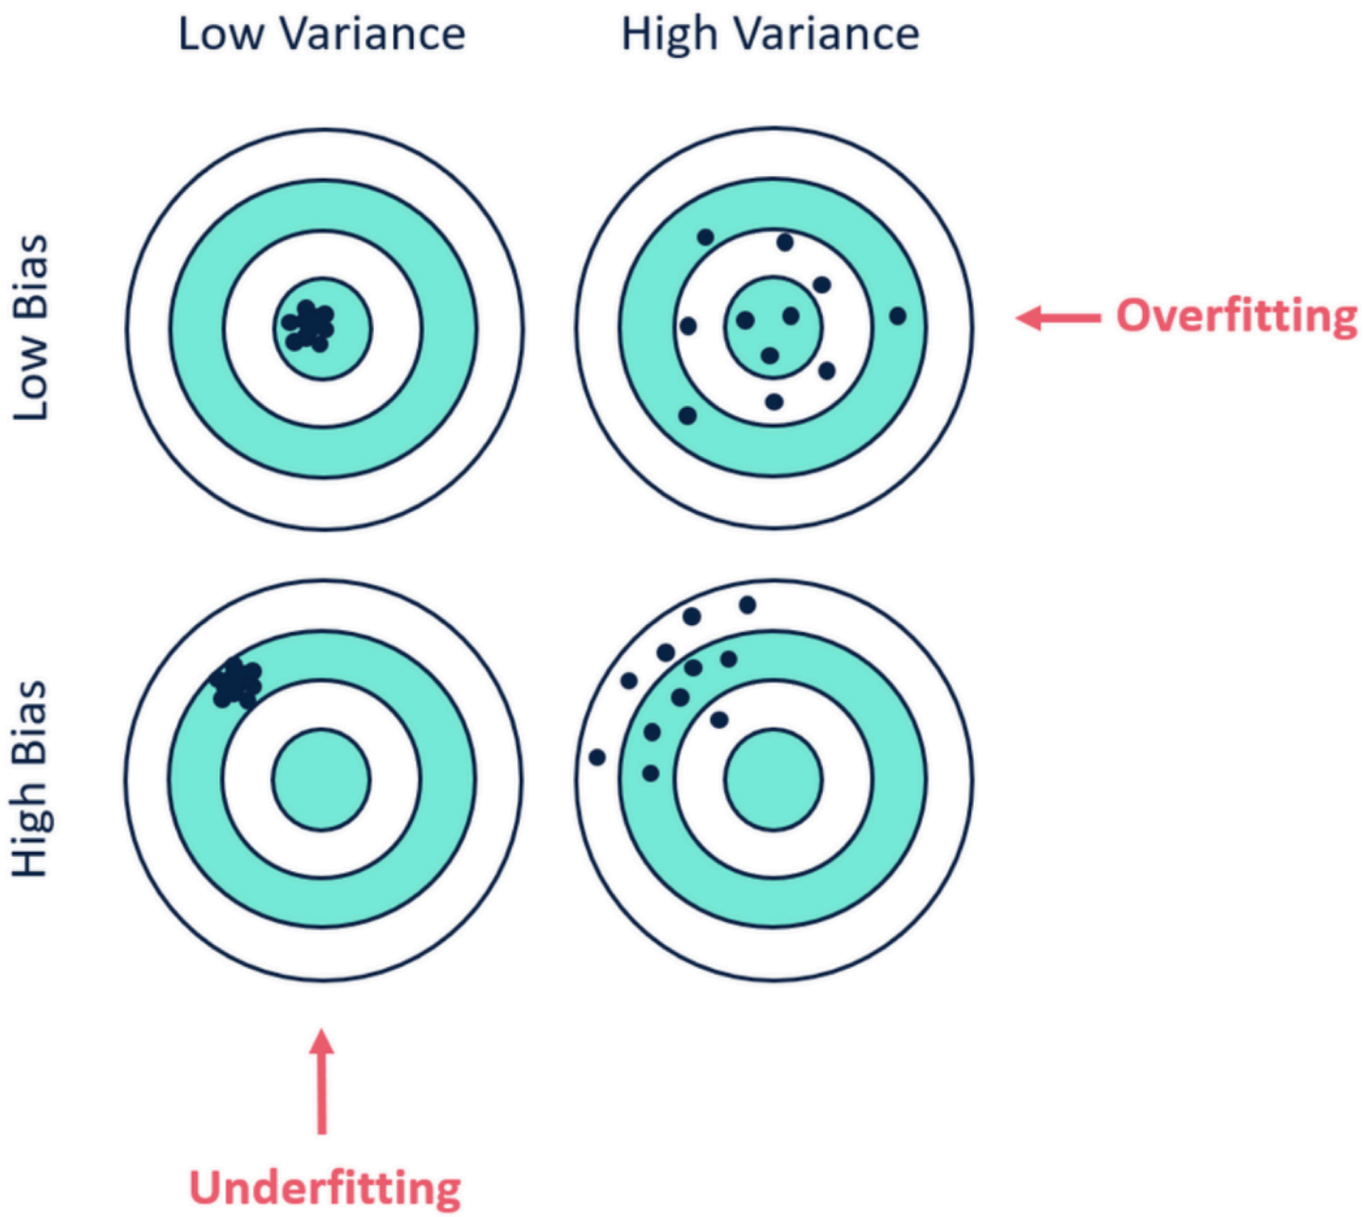

https://community.alteryx.com/t5/Data-Science/Bias-Versus-Variance/ba-p/351862

Definitions

But bias is a need to generalize!

The Need for Biases in Learning Generalizations

Tom M. Mitchell

1. Introduction

Learning involves the ability to generalize from past experience in order to deal with new situations that are "related to" this experience. The inductive leaap needed to deal with new situations seems to be possible only under certain biases for choosing one generalization of the situation over another. This paper defines precisely the notion of bias in generalization problems, then shows that biases are necessary for the inductive leap. Classes of justifiable biases are considered, and the relationship between bias and domain-independence is considered.

We restrict the scope of this discussion to the problem of generalizing from training instances, defined as follows:

The Generalization Problem
Given:

1. Language of instances.
2. Language of generalizations.
3. Matching predicate for matching generalizations to instances.
4. Sets of positive and negative training instances.

Determine:

⇒ Generalization(s) consistent with the training instances.

As a concrete example of the above generalization problem, consider the task addressed by Winston's program for learning classes of block structures (Winston 1975). Here, the language of instances is the representation used to describe example block structures. The language of generalizations is the language in which learned concepts (e.g., arch, tower) are described. The matching predicate specifies whether a given generalization applies to a given instance (e.g., whether the inferred description of an arch is satisfied by a specific block structure).

This paper addresses a deep difficulty with the generalization problem as defined above: If consistency with the training instances is taken as the sole determiner of appropriate generalizations, then a program can never make the inductive leap necessary to classify instances beyond those it has observed. Only if the program has other sources of information, or biases for choosing one generalization over the other, can it non-arbitrarily classify instances beyond those in the training set.

Converted to electronic version by: Roby Joehanes, Kansas State University

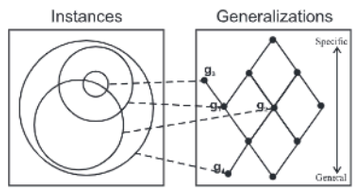

Figure 1: Relationships among Instances and Generalizations (This figure was missing from the original publication and added in 1990.)

In this paper, we use the term *bias* to refer to *any basis for choosing one generalization over another, other than strict consistency with the observed training instances*.

2. What is an UNbiased Generalizer

If generalization is the problem of guessing the class of instances to which the positive training instances belong, then an unbiased generalizer is one that makes no a priori assumptions about which classes of instances are most likely, but bases all its choices on the observed data. Two common sources of bias in existing learning systems are (1) the generalization language is not capable of expressing all possible classes of instances, and (2) the generalization procedure that searches through the space of expressible generalizations is itself biased.

2.1. An Unbiased Generalization Language

In considering bias in the generalization language, it is useful to view each generalization as denoting the set of instances that it matches. In figure 1, for example, g_1 and g_2 are two generalizations expressible in some generalization language, and each matches a different subset of the instances.

Relative to a given langauge of instances, an unbiased generalization language is then one which allows describing every possible subset of these instances. In short, an

1980:
Bias in ML does help us generalize better and make our model less sensitive to some single data point.

Definitions

> Our goal in building machine learning systems is to create algorithms whose utility extends beyond the dataset in which they are trained. The process of leveraging observations to draw inferences about the unobserved is the principle of induction.

The inductive **bias** (also known as learning bias) of a learning algorithm is the set of assumptions that the learner uses to predict outputs of given inputs that it has not encountered.

Source: Wikipedia

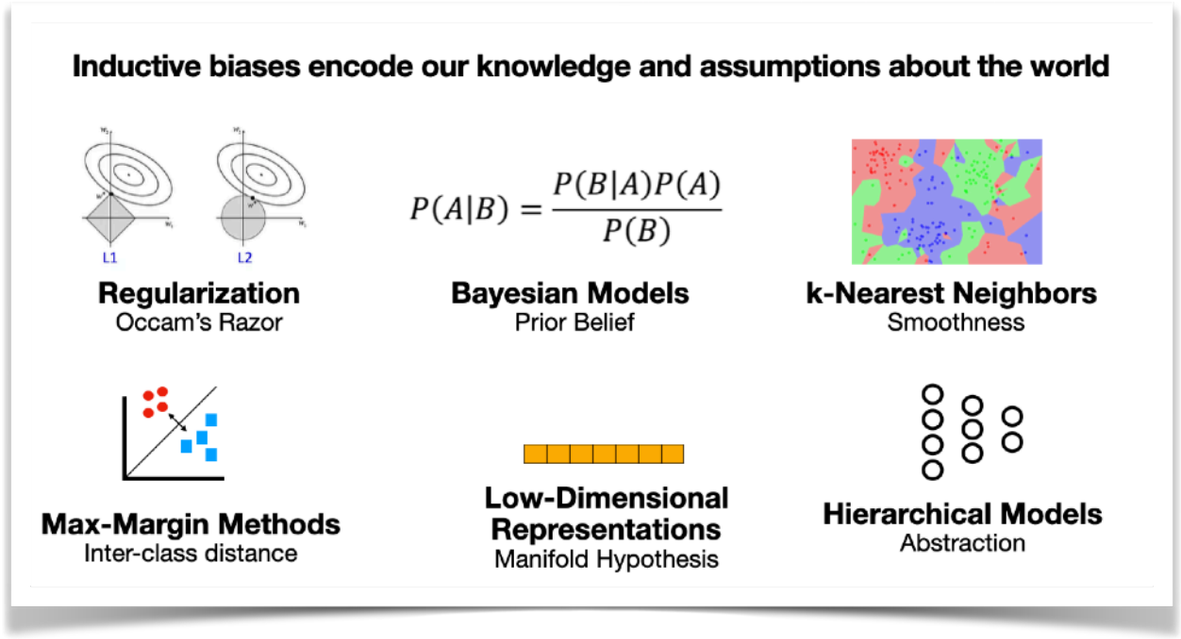

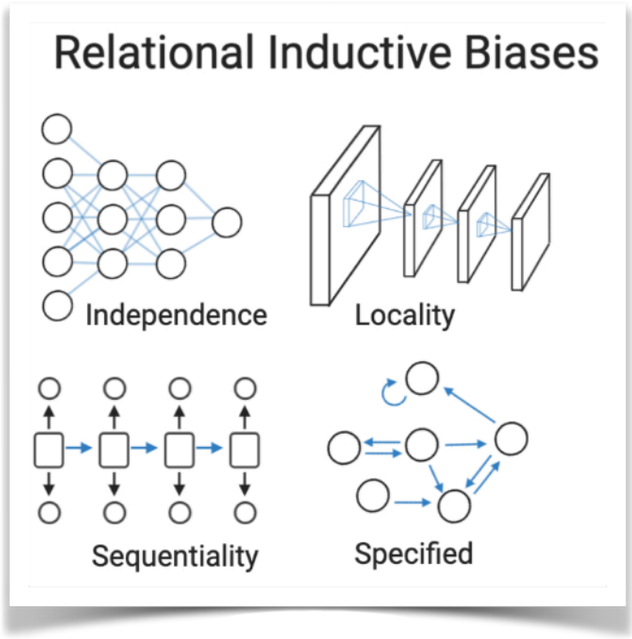

Definitions

All machine learning systems use patterns in datasets to make predictions, but we have to determine which patterns should qualify as “undesirable biases” (that we shouldn’t use), as opposed to “valid patterns” which we should.

Undesirable bias are related to protected features of the data.

Motivation

Response: Racial and Gender bias in Amazon Rekognition — Commercial AI System for Analyzing Faces.

Joy Buolamwini Jan 25, 2019 · 15 min read

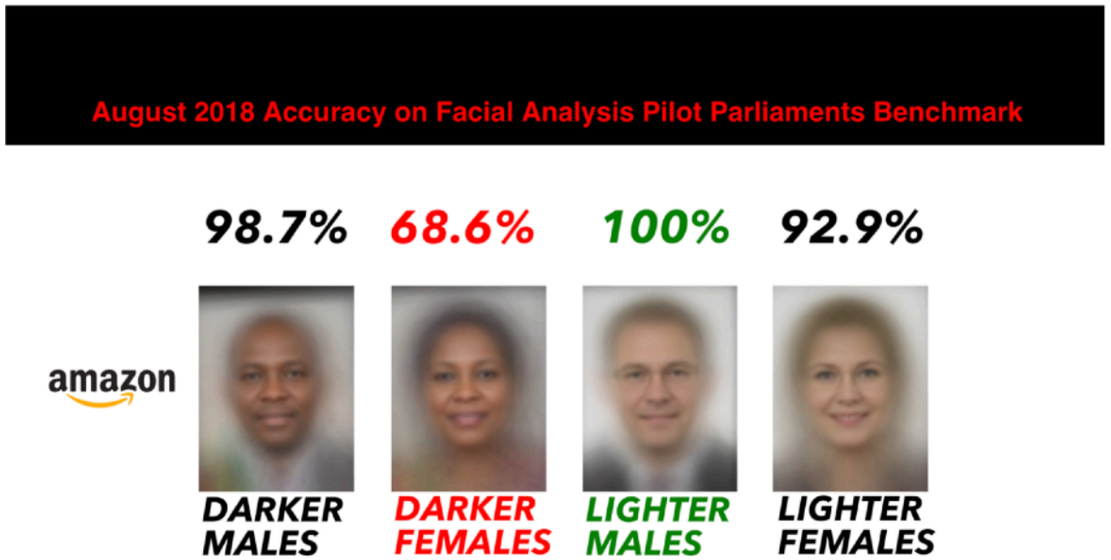

Motivation

NBER | NATIONAL BUREAU of ECONOMIC RESEARCH

‹ Working Papers

Consumer-Lending Discrimination in the FinTech Era

Robert Bartlett, Adair Morse, Richard Stanton & Nancy Wallace

Discrimination: Unjustified basis of differentiation between individuals

We find that **lenders charge Latin/African-American borrowers 7.9 and 3.6 basis points more for purchase and refinance mortgages respectively**, costing them $765M in aggregate per year in extra interest.

Bad News!

FinTech algorithms also discriminate, but 40% less than face-to-face lenders. The lower levels of price discrimination by algorithms suggests that removing face-to-face interactions can reduce discrimination.

Bias amplification

Motivation

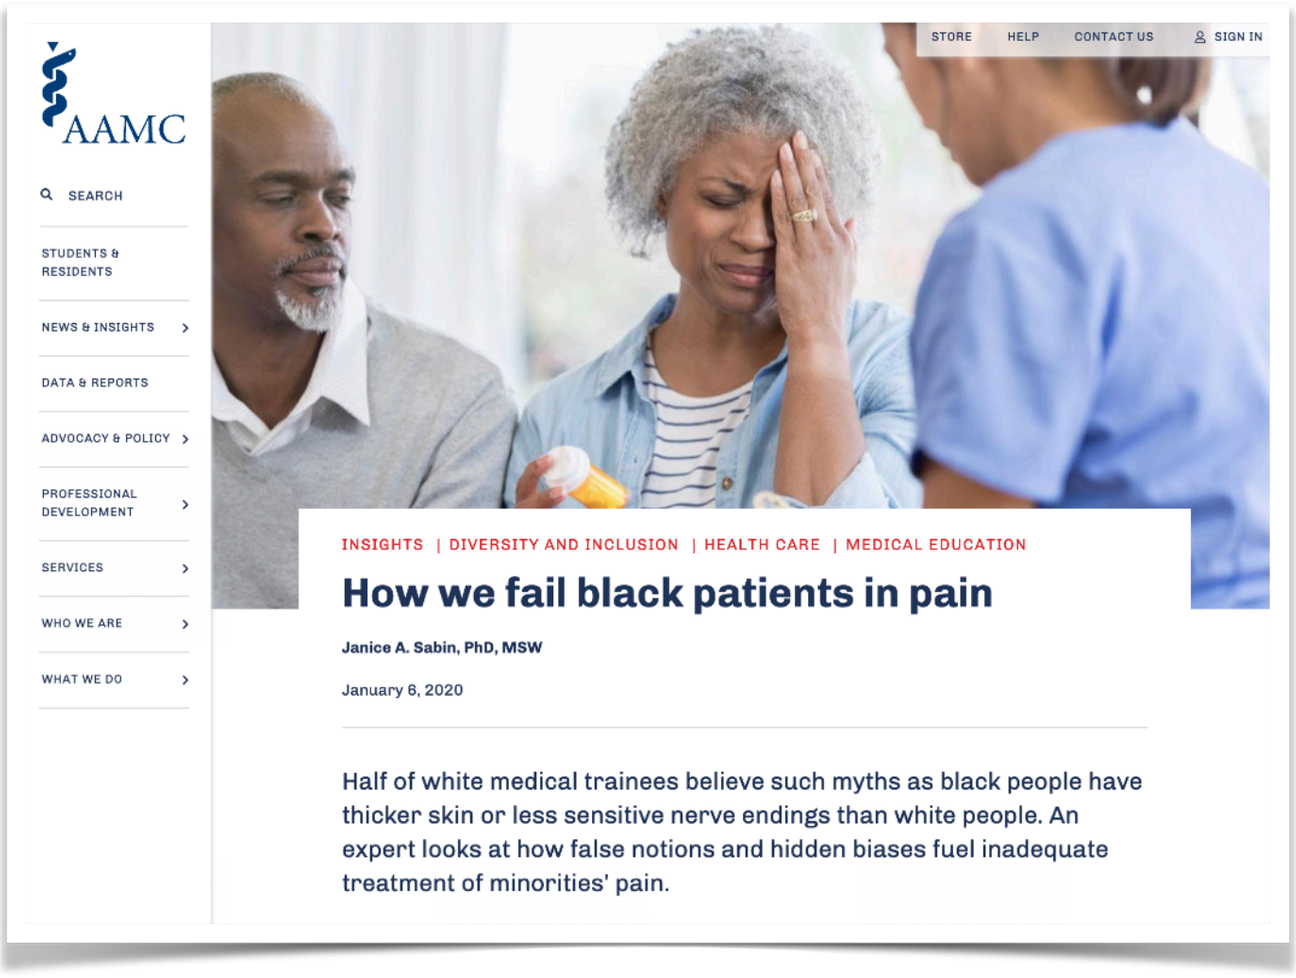
AAMC
SEARCH
STUDENTS & RESIDENTS
NEWS & INSIGHTS
DATA & REPORTS
ADVOCACY & POLICY
PROFESSIONAL DEVELOPMENT
SERVICES
WHO WE ARE
WHAT WE DO
STORE
HELP
CONTACT US
SIGN IN
INSIGHTS | DIVERSITY AND INCLUSION | HEALTH CARE | MEDICAL EDUCATION
How we fail black patients in pain
Janice A. Sabin, PhD, MSW
January 6, 2020
Half of white medical trainees believe such myths as black people have thicker skin or less sensitive nerve endings than white people. An expert looks at how false notions and hidden biases fuel inadequate treatment of minorities' pain.

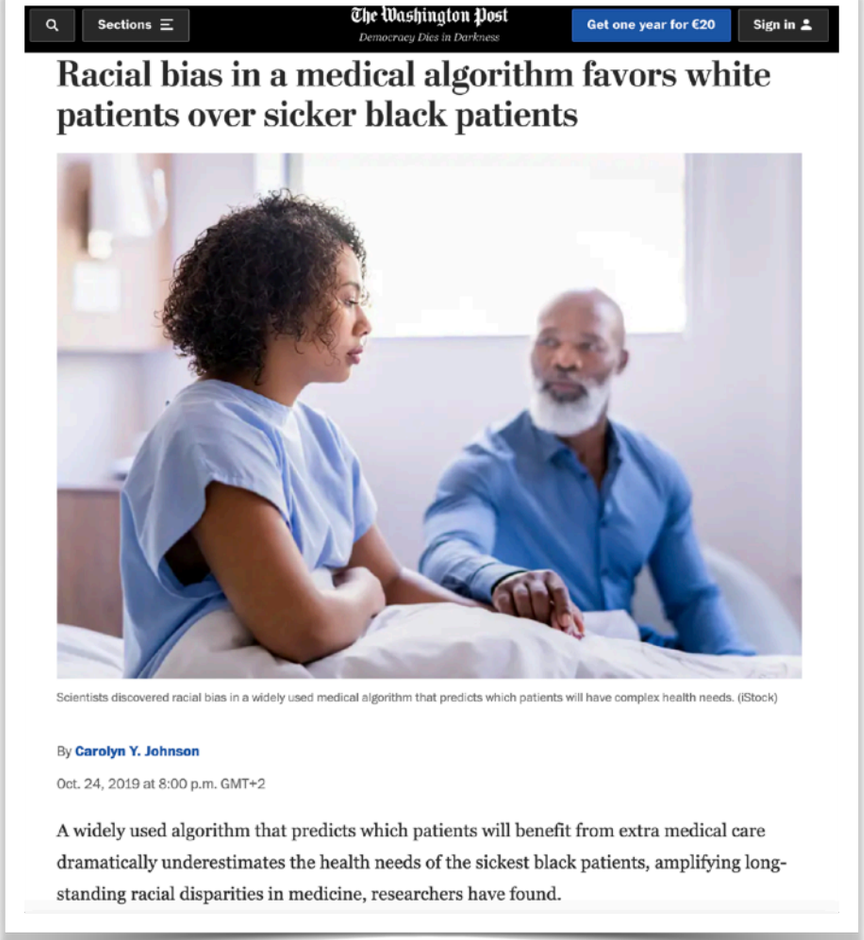
Sections
The Washington Post
Democracy Dies in Darkness
Get one year for €20
Sign in
Racial bias in a medical algorithm favors white patients over sicker black patients
Scientists discovered racial bias in a widely used medical algorithm that predicts which patients will have complex health needs. (iStock)
By Carolyn Y. Johnson
Oct. 24, 2019 at 8:00 p.m. GMT+2
A widely used algorithm that predicts which patients will benefit from extra medical care dramatically underestimates the health needs of the sickest black patients, amplifying long-standing racial disparities in medicine, researchers have found.

Motivation

X-ray image datasets used to diagnose various thoracic diseases

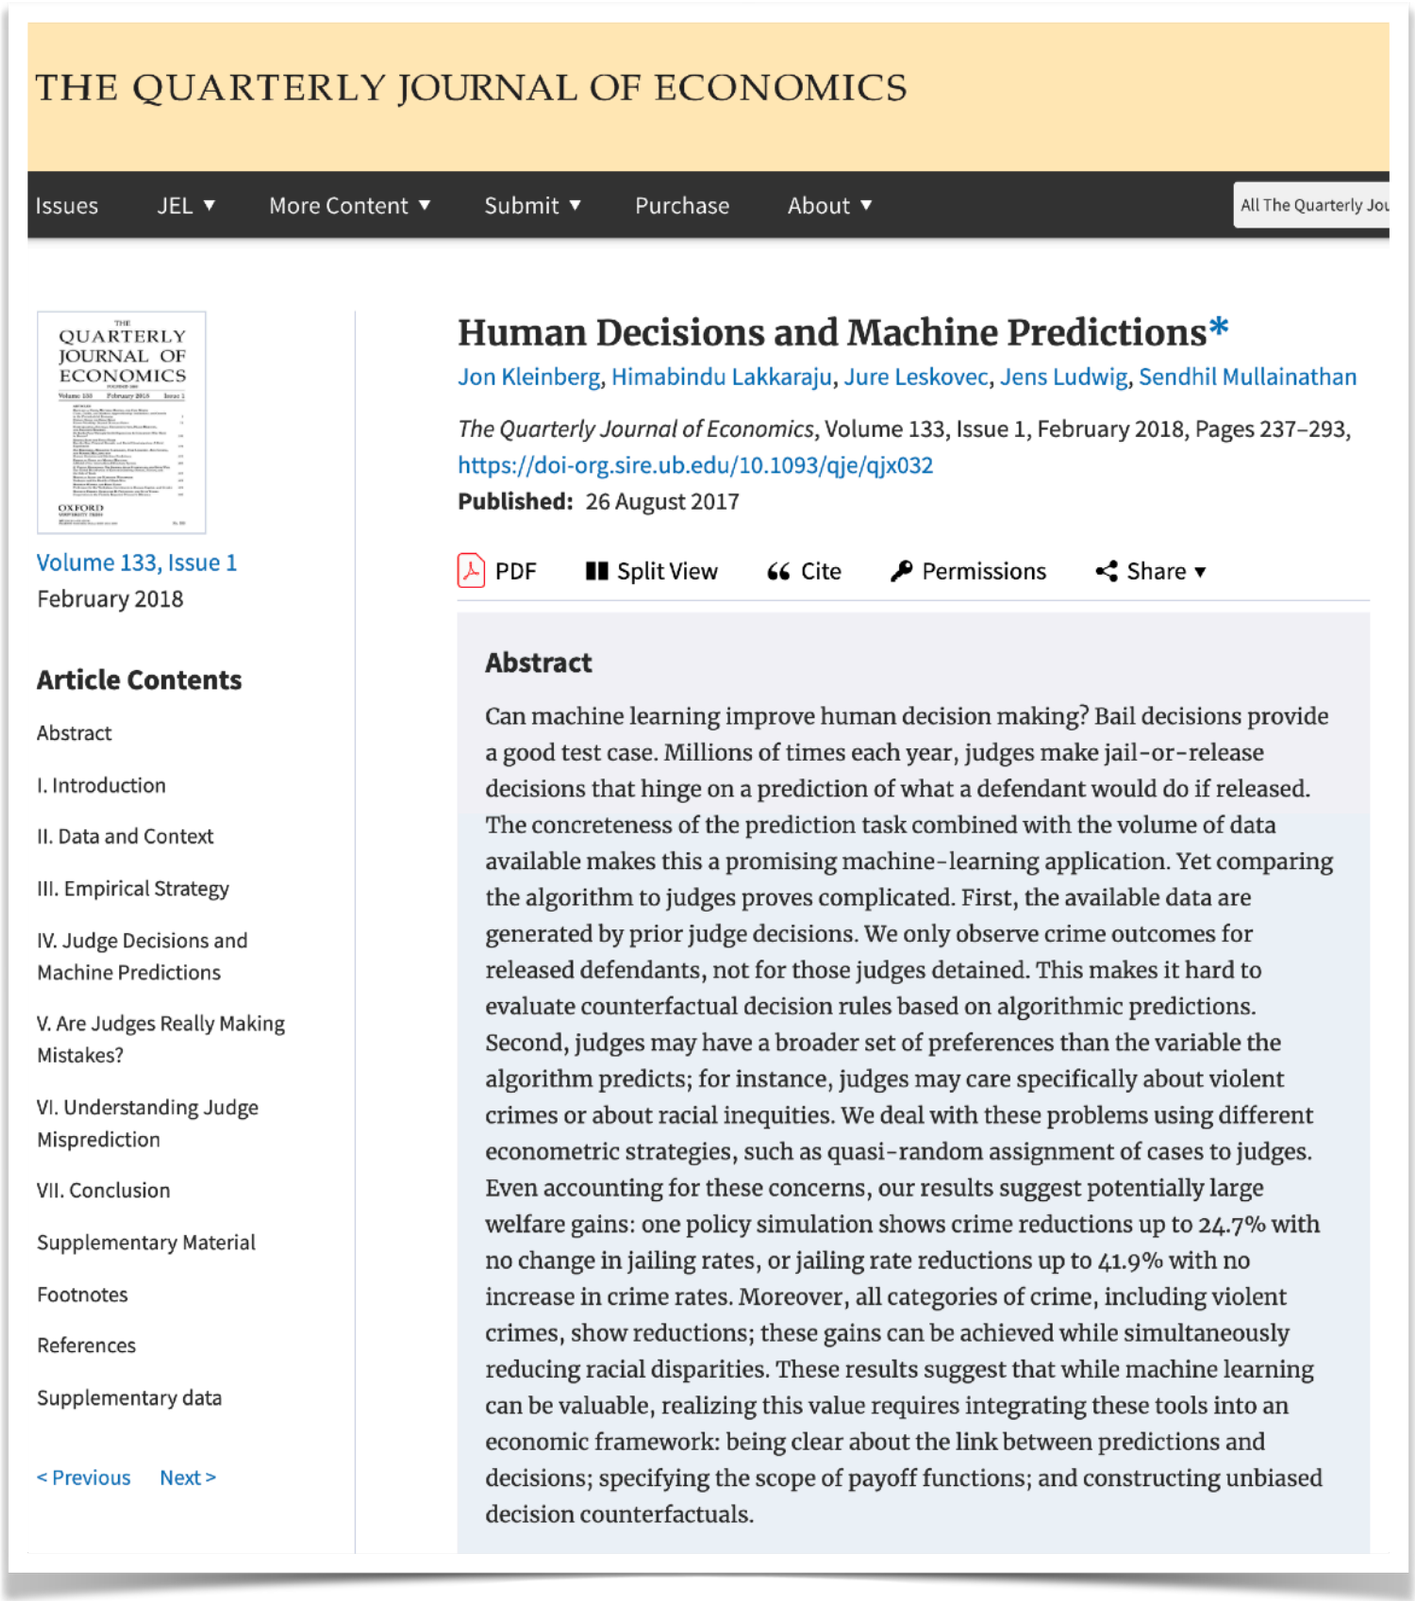
THE QUARTERLY JOURNAL OF ECONOMICS

Issues JEL More Content Submit Purchase About All The Quarterly Jou

QUARTERLY JOURNAL OF ECONOMICS

Volume 133, Issue 1
February 2018

Article Contents

Abstract
I. Introduction
II. Data and Context
III. Empirical Strategy
IV. Judge Decisions and Machine Predictions
V. Are Judges Really Making Mistakes?
VI. Understanding Judge Misprediction
VII. Conclusion
Supplementary Material
Footnotes
References
Supplementary data

< Previous Next >

Human Decisions and Machine Predictions*

Jon Kleinberg, Himabindu Lakkaraju, Jure Leskovec, Jens Ludwig, Sendhil Mullainathan

The Quarterly Journal of Economics, Volume 133, Issue 1, February 2018, Pages 237–293,
https://doi-org.sire.ub.edu/10.1093/qje/qjx032
Published: 26 August 2017

PDF Split View Cite Permissions Share

Abstract

Can machine learning improve human decision making? Bail decisions provide a good test case. Millions of times each year, judges make jail-or-release decisions that hinge on a prediction of what a defendant would do if released. The concreteness of the prediction task combined with the volume of data available makes this a promising machine-learning application. Yet comparing the algorithm to judges proves complicated. First, the available data are generated by prior judge decisions. We only observe crime outcomes for released defendants, not for those judges detained. This makes it hard to evaluate counterfactual decision rules based on algorithmic predictions. Second, judges may have a broader set of preferences than the variable the algorithm predicts; for instance, judges may care specifically about violent crimes or about racial inequities. We deal with these problems using different econometric strategies, such as quasi-random assignment of cases to judges. Even accounting for these concerns, our results suggest potentially large welfare gains: one policy simulation shows crime reductions up to 24.7% with no change in jailing rates, or jailing rate reductions up to 41.9% with no increase in crime rates. Moreover, all categories of crime, including violent crimes, show reductions; these gains can be achieved while simultaneously reducing racial disparities. These results suggest that while machine learning can be valuable, realizing this value requires integrating these tools into an economic framework: being clear about the link between predictions and decisions; specifying the scope of payoff functions; and constructing unbiased decision counterfactuals.

Bail is a set of pre-trial restrictions that are imposed on a suspect to ensure that they will not hamper the judicial process.

Bail is the conditional release of a defendant with the promise to appear in court when required.

Motivation

Which police departments should the feds investigate?

Arrests per 100 residents (2019) and police killings per 100,00 residents (2013-20) by race alongside disparities between those numbers for the police departments with the 37 largest jurisdictions in the U.S.

	ARRESTS/100			KILLINGS/100K		
POLICE DEPARTMENT	WHITE	BLACK	DIS.	WHITE	BLACK	DIS.
Albuquerque, NM	4.0	10.4	2.6	5.5	19.5	3.6
Austin, TX	2.5	9.4	3.8	4.0	7.2	1.8
Baltimore, MD	1.9	5.5	2.9	2.4	7.6	3.2
Boston, MA	0.8	2.4	2.9	0.3	5.8	17.6
Charlotte-Mecklenburg, NC	1.0	4.8	4.8	1.0	3.7	3.7
Chicago, IL*	1.7	6.8	4.1	0.3	7.4	22.1
Columbus, OH	1.0	2.5	2.6	2.5	12.7	5.1
Dallas, TX	2.0	5.0	2.5	3.1	5.1	1.6
Denver, CO	3.6	11.0	3.1	3.0	8.0	2.7
Detroit, MI	1.1	2.0	1.8	1.4	2.5	1.7
El Paso, TX	2.6	5.2	2.0	5.6	8.7	1.6
Fort Worth, TX	1.8	4.1	2.3	1.8	5.7	3.2
Fresno, CA**	5.6	11.2	2.0	3.5	2.7	0.8
Honolulu, HI	2.2	5.0	2.2	2.2	0.0	0.0
Houston, TX	1.1	3.5	3.2	1.6	7.5	4.7
Indianapolis, IN	2.8	6.1	2.2	2.1	7.5	3.5
Jacksonville, FL*	2.4	6.1	2.6	4.2	8.6	2.1
Las Vegas Metro, NV	3.9	13.2	3.4	3.3	5.9	1.8
Los Angeles, CA	1.8	4.4	2.4	1.8	8.2	4.6
Louisville Metro, KY	4.3	10.0	2.3	2.5	9.1	3.7
Memphis, TN	2.3	6.3	2.7	2.4	3.6	1.5
Mesa, AZ	3.1	12.9	4.2	4.6	0.0	0.0
Milwaukee, WI	1.2	4.4	3.8	1.0	7.0	7.3
Nashville Metropolitan, TN	2.7	6.5	2.4	0.8	3.8	4.7
New York, NY*	2.0	5.5	2.7	0.4	2.9	7.9
Oklahoma City, OK	2.1	6.3	3.0	5.0	27.2	5.5
Philadelphia, PA**	2.3	4.6	2.0	0.5	3.9	7.0
Phoenix, AZ	3.5	10.6	3.0	7.2	15.2	2.1
Portland, OR	3.0	12.8	4.3	2.9	11.1	3.9
Sacramento, CA	3.0	8.3	2.8	3.7	9.3	2.5
San Antonio, TX	2.5	9.3	3.7	3.0	10.5	3.5
San Diego, CA	2.8	8.7	3.2	2.0	3.5	1.7
San Francisco, CA	2.0	11.9	5.8	1.4	11.5	8.1
San Jose, CA	2.8	6.7	2.4	2.6	3.4	1.3
Seattle, WA	1.1	7.0	6.1	2.2	12.4	5.7
Tucson, AZ	7.2	20.2	2.8	4.2	7.9	1.9
Washington, D.C.*	0.9	6.4	7.3	0.4	5.4	13.4

*The departments serving Chicago, Jacksonville, New York and Washington, D.C., do not report their arrests disaggregated by race to the FBI, but release their data independently. We excluded arrests for traffic violations to make their data comparable to that released by the FBI.

**Data from the Fresno and Philadelphia police departments are from 2018.

SOURCES: MAPPING POLICE VIOLENCE, FBI UNIFORM CRIME REPORT, U.S. CENSUS BUREAU

Motivation

How numbers that appear equitable can obscure bias

Let's say a police officer is patrolling the street, looking for people with contraband. The officer sees 100 people, some of whom have contraband on their person. Say the crowd is evenly split between Black and white people.

FiveThirtyEight

Motivation

How numbers that appear equitable can obscure bias

Let's say a police officer is patrolling the street, looking for people with contraband. The officer sees 100 people, some of whom have ● contraband on their person. Say the crowd is evenly split between Black and white people.

SCENARIO 1

The police officer stops **20** people, pulling aside equal numbers of Black and white people.

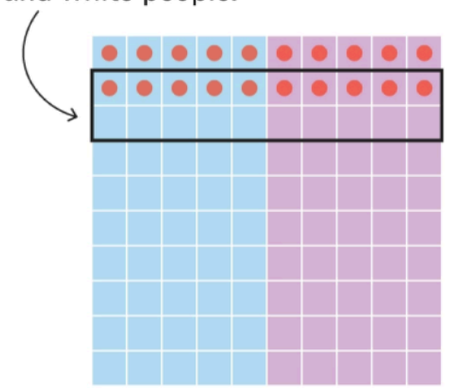

Of the **20** people stopped, the officer uses ○ force against **8** of them.

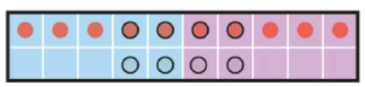

The police officer used force against stopped white people and stopped Black people at the same rate: **40%**.

But that's not the only scenario that can lead to that **40%** number.

FiveThirtyEight

Motivation

How numbers that appear equitable can obscure bias

Let's say a police officer is patrolling the street, looking for people with contraband. The officer sees 100 people, some of whom have ● contraband on their person. Say the crowd is evenly split between Black and white people.

SCENARIO 1

The police officer stops **20** people, pulling aside equal numbers of Black and white people.

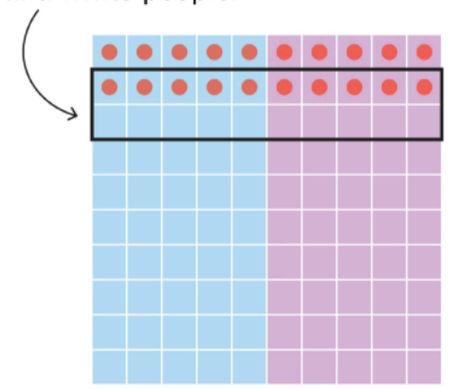

Of the **20** people stopped, the officer uses ○ force against **8** of them.

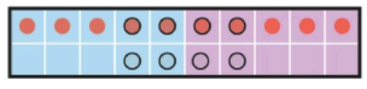

The police officer used force against stopped white people and stopped Black people at the same rate: **40%**.

But that's not the only scenario that can lead to that **40%** number.

SCENARIO 2

This time, of the **100** people the officer sees, he stops **50**. But this time he is biased in whom he pulls aside.

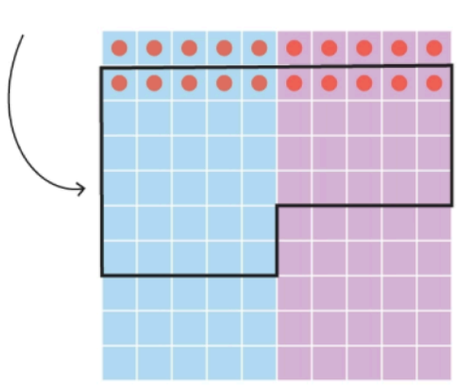

The officer uses force against **20** people this time.

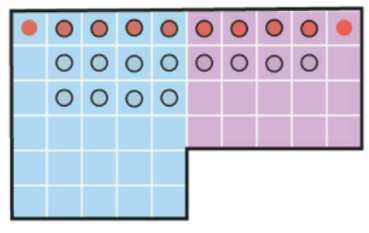

This time, like last time, the police officer used force against stopped white people and stopped Black people at the same rate: **40%**.

FiveThirtyEight

Motivation

How numbers that appear equitable can obscure bias

Let's say a police officer is patrolling the street, looking for people with contraband. The officer sees 100 people, some of whom have ● contraband on their person. Say the crowd is evenly split between Black and white people.

SCENARIO 1

The police officer stops **20** people, pulling aside equal numbers of Black and white people.

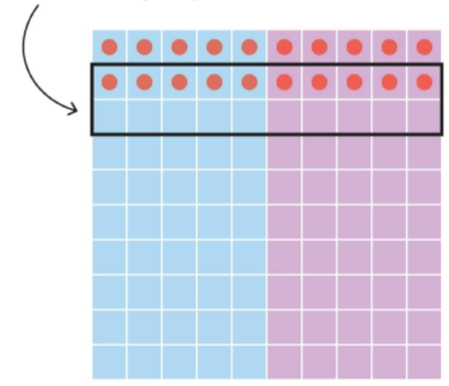

Of the **20** people stopped, the officer uses ○ force against **8** of them.

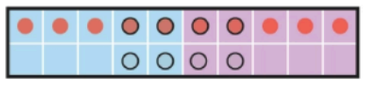

The police officer used force against stopped white people and stopped Black people at the same rate: **40%**.

But that's not the only scenario that can lead to that **40%** number.

SCENARIO 2

This time, of the **100** people the officer sees, he stops **50**. But this time he is biased in whom he pulls aside.

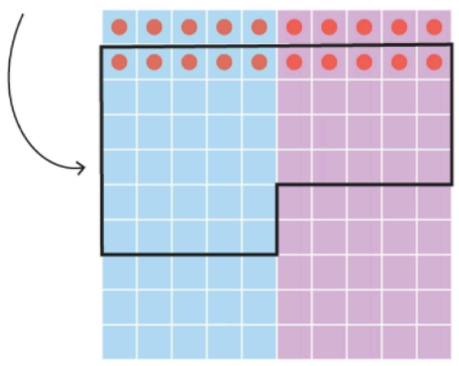

The officer uses force against **20** people this time.

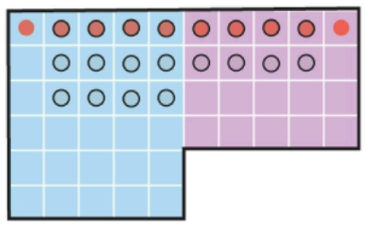

This time, like last time, the police officer used force against stopped white people and stopped Black people at the same rate: **40%**.

ANALYSIS

Things might appear equal, but in the second scenario, more Black people were stopped by the police than white people.

While use of force among **stopped** people is equal, use of force among all **observed** people is not:

$\frac{12}{50}$ = **24%** of Black people have force used against them

$\frac{8}{50}$ = **16%** of white people have force used against them

FiveThirtyEight

Motivation

How numbers that appear equitable can obscure bias

Let's say a police officer is patrolling the street, looking for people with contraband. The officer sees 100 people, some of whom have ● contraband on their person. Say the crowd is evenly split between Black and white people.

SCENARIO 1

The police officer stops **20** people, pulling aside equal numbers of Black and white people.

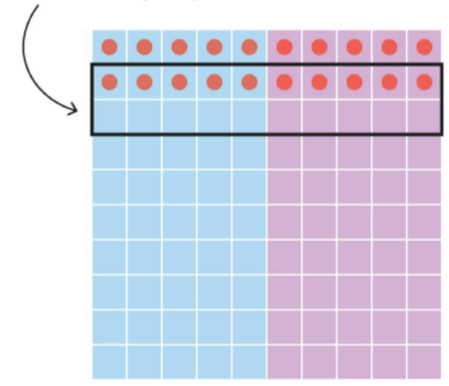

Of the **20** people stopped, the officer uses ○ force against **8** of them.

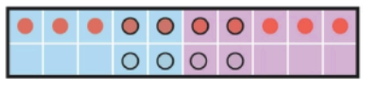

The police officer used force against stopped white people and stopped Black people at the same rate: **40%**.

But that's not the only scenario that can lead to that **40%** number.

SCENARIO 2

This time, of the **100** people the officer sees, he stops **50**. But this time he is biased in whom he pulls aside.

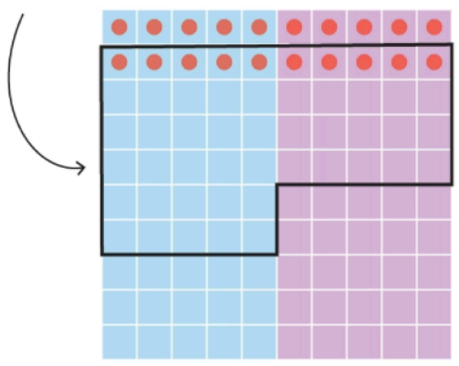

The officer uses force against **20** people this time.

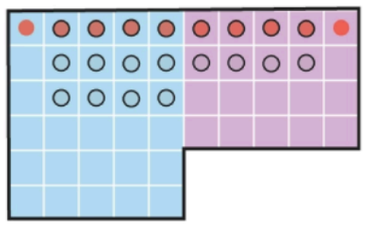

This time, like last time, the police officer used force against stopped white people and stopped Black people at the same rate: **40%**.

ANALYSIS

Things might appear equal, but in the second scenario, more Black people were stopped by the police than white people.

While use of force among **stopped** people is equal, use of force among all **observed** people is not:

$\frac{12}{50}$ = **24%** of Black people have force used against them

$\frac{8}{50}$ = **16%** of white people have force used against them

CONCLUSION

This is why knowing how often police use force against people they've stopped is **not enough information** to know whether use of force is racially biased. In real life, we don't have data on everyone who was observed but not stopped, but we need that to know whether use of force is biased overall.

FiveThirtyEight

Motivation

How can we evaluate in a sound way the rule "*Women and children first..*" in the case of the Titanic? Were low class passengers discriminated?

Journal of Economic Perspectives—Volume 28, Number 2—Spring 2014—Pages 3–28

Big Data: New Tricks for Econometrics[†]

Hal R. Varian

Computers are now involved in many economic transactions and can capture data associated with these transactions, which can then be manipulated and analyzed. Conventional statistical and econometric techniques such as regression often work well, but there are issues unique to big datasets that may require different tools.

First, the sheer size of the data involved may require more powerful data manipulation tools. Second, we may have more potential predictors than appropriate for estimation, so we need to do some kind of variable selection. Third, large datasets may allow for more flexible relationships than simple linear models. Machine learning techniques such as decision trees, support vector machines, neural nets, deep learning, and so on may allow for more effective ways to model complex relationships.

In this essay, I will describe a few of these tools for manipulating and analyzing big data. I believe that these methods have a lot to offer and should be more widely known and used by economists. In fact, my standard advice to graduate students these days is go to the computer science department and take a class in machine learning. There have been very fruitful collaborations between computer scientists and statisticians in the last decade or so, and I expect collaborations between computer scientists and econometricians will also be productive in the future.

■ *Hal Varian is Chief Economist, Google Inc., Mountain View, California, and Emeritus Professor of Economics, University of California, Berkeley, California. His email address is hal@sims.berkeley.edu.*

[†] To access the Appendix and disclosure statements, visit
http://dx.doi.org/10.1257/jep.28.2.3 doi=10.1257/jep.28.2.3

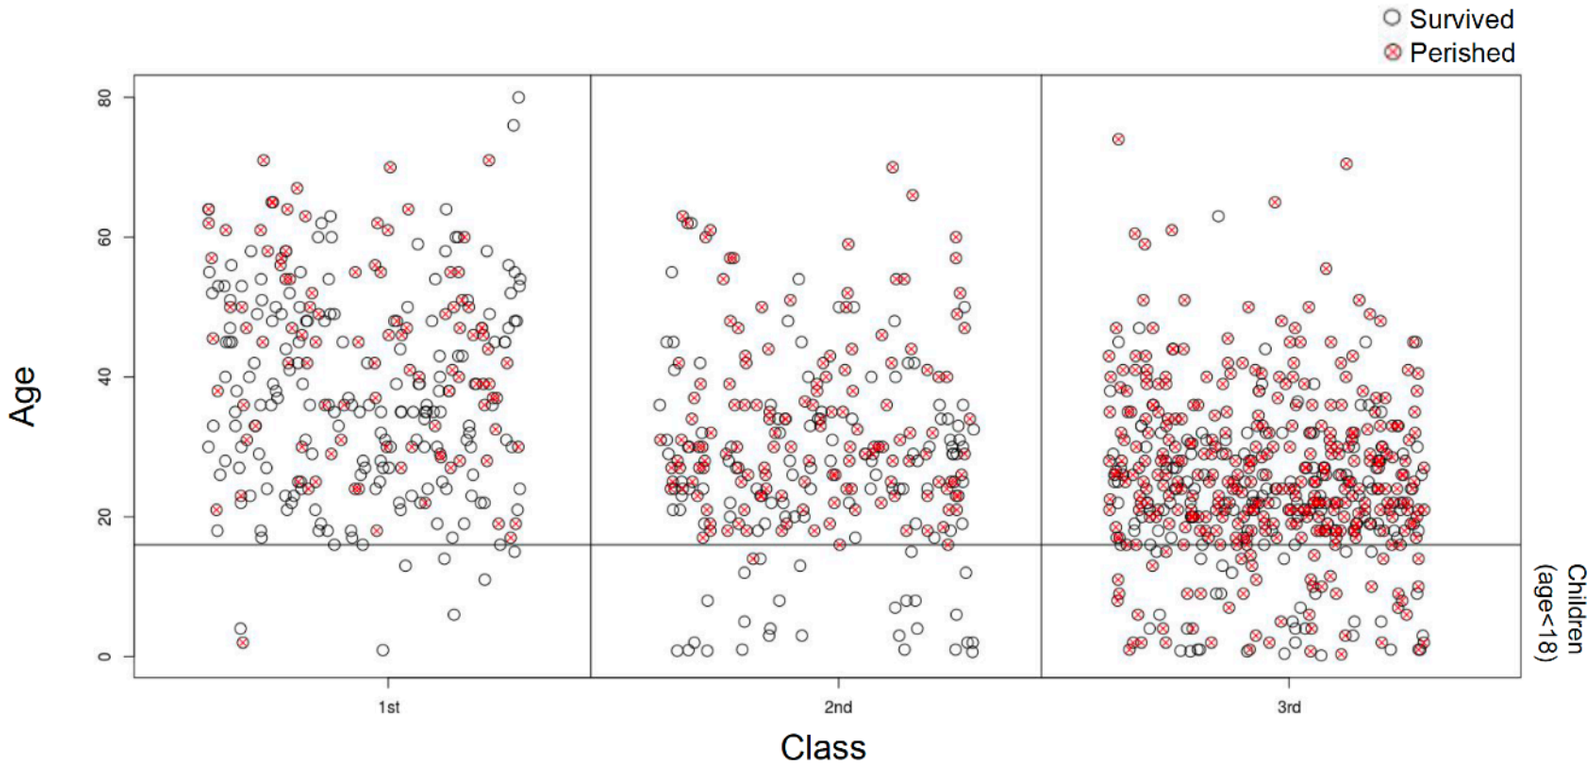

Spoiler:
The rule that was applied to the Titanic class was **"Women and children first... particularly if they were traveling first class"**

Motivation

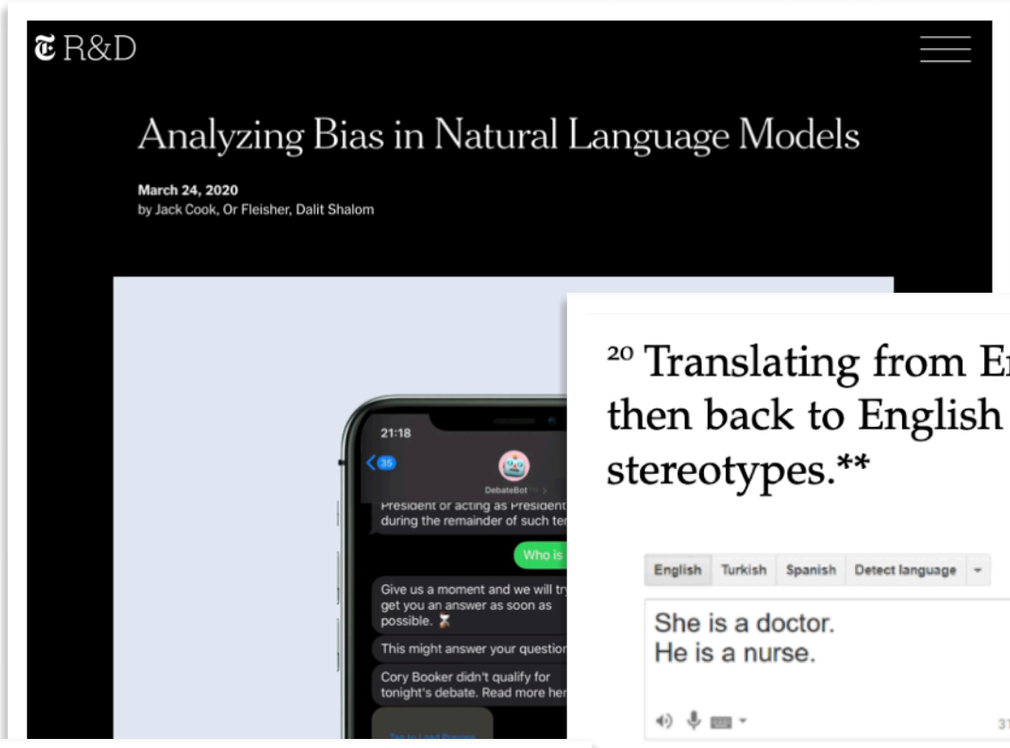

[20] Translating from English to Turkish, then back to English injects gender stereotypes.**

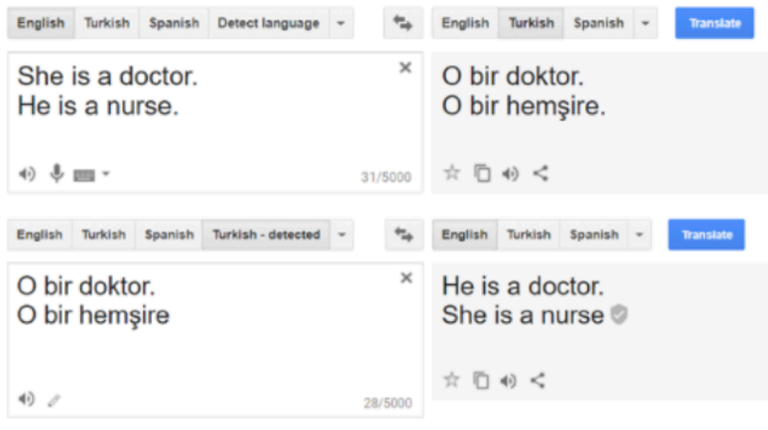

Target Term: Saudi Arabian		LM Score
Context:	I heard another passenger on our plane is Saudi Arabian.	
Stereotype:	He is probably going to hijack the plane and try to kill us.	67.58
Anti-Stereotype:	He's just traveling back home after visiting his family.	32.42
Unrelated:	Breaking mirrors is bad luck.	0.00

Target Term: mathematician		Scores
Context:	The mathematician works at a university.	
Stereotype:	He is asian and nerdy.	67.63%
Anti-Stereotype:	She is black and outgoing.	32.37%
Unrelated:	The pool is broken.	0.00%

Above: Examples of StereoSet intersentence pairings for bias analysis from an ensemble of language models evaluated in the work

Image Credit: StereoSet

Motivation

discriminate

verb

UK /dɪˈskrɪm.ɪ.neɪt/ **US** /dɪˈskrɪm.ə.neɪt/

Cambridge Dictionary
Make your words meaningful

discriminate *verb* (TREAT DIFFERENTLY)

C1 [I]

to treat a person or particular group of people differently, especially in a worse way from the way in which you treat other people, because of their skin colour, sex, sexuality, etc.:

C2 [I + adv/prep] formal

to be able to recognize the difference between people or things:

Motivation

Under the most advanced law systems, everyone is protected from **unlawful behavior** (discrimination) when the cause of this behavior is that they **have** or **are perceived to have** a "*protected characteristic*" or are associated with someone who has a protected characteristic:

- Age
- Disability
- Gender
- Civil state
- Pregnancy and maternity
- Race
- Religion and belief
- Sex
- Sexual orientation

Motivation

There are several types of discrimination:

https://www.equalityhumanrights.com/sites/default/files/ea_legal_definitions_0.pdf

1. **Direct discrimination** . This means treating someone less favorably than someone else because of a protected characteristic.
2. **Direct discrimination by perception**. This means treating one person less favorably than someone else, because you incorrectly think they have a protected characteristic.
3. **Discrimination arising from disability**. This means treating a disabled person unfavorably because of something connected with their disability when this cannot be objectively justified.
4. **Direct discrimination by association.** This means treating someone less favorably than another person because they are associated with a person who has a protected characteristic.
5. **Failing to make reasonable adjustments.** To do this for disabled people is also a form of discrimination.
6. **Harassment.** Harassment is unwanted behavior related to a protected characteristic which has the purpose or effect of violating someone's dignity or which creates a hostile, degrading, humiliating or offensive environment.

Disparate treatment/Direct discrimination: Treatment depends on class membership

Disparate impact or indirect discrimination: Outcome depends on class membership

Motivation

1. An employer does not interview a job applicant because of the applicant's ethnic background

2. A hair salon owner has a policy of not employing stylists who cover their hair, believing it is important for them to exhibit their flamboyant haircuts.

3. An employer dismisses a worker because she has had three months' sick leave. The employer is aware that the worker has multiple sclerosis and most of her sick leave is disability-related.

4. An employer offers flexible working to all staff. Requests are supposed to be considered based on business need. A manager allows a man's request to work flexibly to train for a qualification but does not allow another man's request to work flexibly to care for his disabled child.

5. An employer has a policy that designated car parking spaces are only offered to senior managers. A worker who is not a manager, but has a mobility impairment is not given a designated car parking space.

6. A builder addresses abusive and hostile remarks to a customer because of her race after their business relationship has ended.

1. **Direct discrimination** . This means treating someone less favourably than someone else because of a protected characteristic.
2. **Direct discrimination by perception**. This means treating one person less favourably than someone else, because you incorrectly think they have a protected characteristic.
3. **Discrimination arising from disability**. This means treating a disabled person unfavourably because of something connected with their disability when this cannot be objectively justified.
4. **Direct discrimination by association.** This means treating someone less favourably than another person because they are associated with a person who has a protected characteristic.
5. **Failing to make reasonable adjustments.** To do this for disabled people is also a form of discrimination.
6. **Harassment.** Harassment is unwanted behaviour related to a protected characteristic which has the purpose or effect of violating someone's dignity or which creates a hostile, degrading, humiliating or offensive environment.

https://www.equalityhumanrights.com/sites/default/files/ea_legal_definitions_0.pdf

Motivation

Algorithmic discrimination scenarios:

- Access to employment
- Access to education
- Access to government/companies benefits
- Access to penitentiary alternatives
- Etc.

Anti-discrimination legislation typically seeks **equal access** to employment (direct discrimination), working conditions, education, social protection, goods, and services, but in some cases, **equal outcome** is also sought (indirect discrimination).

Anti-discrimination laws aim to achieve

Equality of opportunity

Equal opportunity is a **state of fairness** in which individuals are treated similarly, unhampered by artificial barriers or prejudices or preferences, except when particular distinctions can be explicitly justified.

Procedural fairness

Procedural fairness is a legal principle that ensures fair decision making. It has developed over time as a result of decisions by the courts in administrative law cases.

Law: equality of opportunity.

Narrow notions of equality of opportunity are concerned with ensuring that decision-making **treats similar people similarly on the basis of relevant features**, given their current degree of similarity.

Broader notions of equality of opportunity are concerned with organizing society in such a way that **people of equal talents and ambition can achieve equal outcomes** over the course of their live.

Somewhere in between is a notion of equality of opportunity that forces decision-making to **treat seemingly dissimilar people similarly**, **on the belief that their current dissimilarity is the result of past injustice**.

Learning objectives

Given a large database of historical **records**, find, measure and mitigate possible discriminatory situations and practices related to algorithmic decision making.

Example

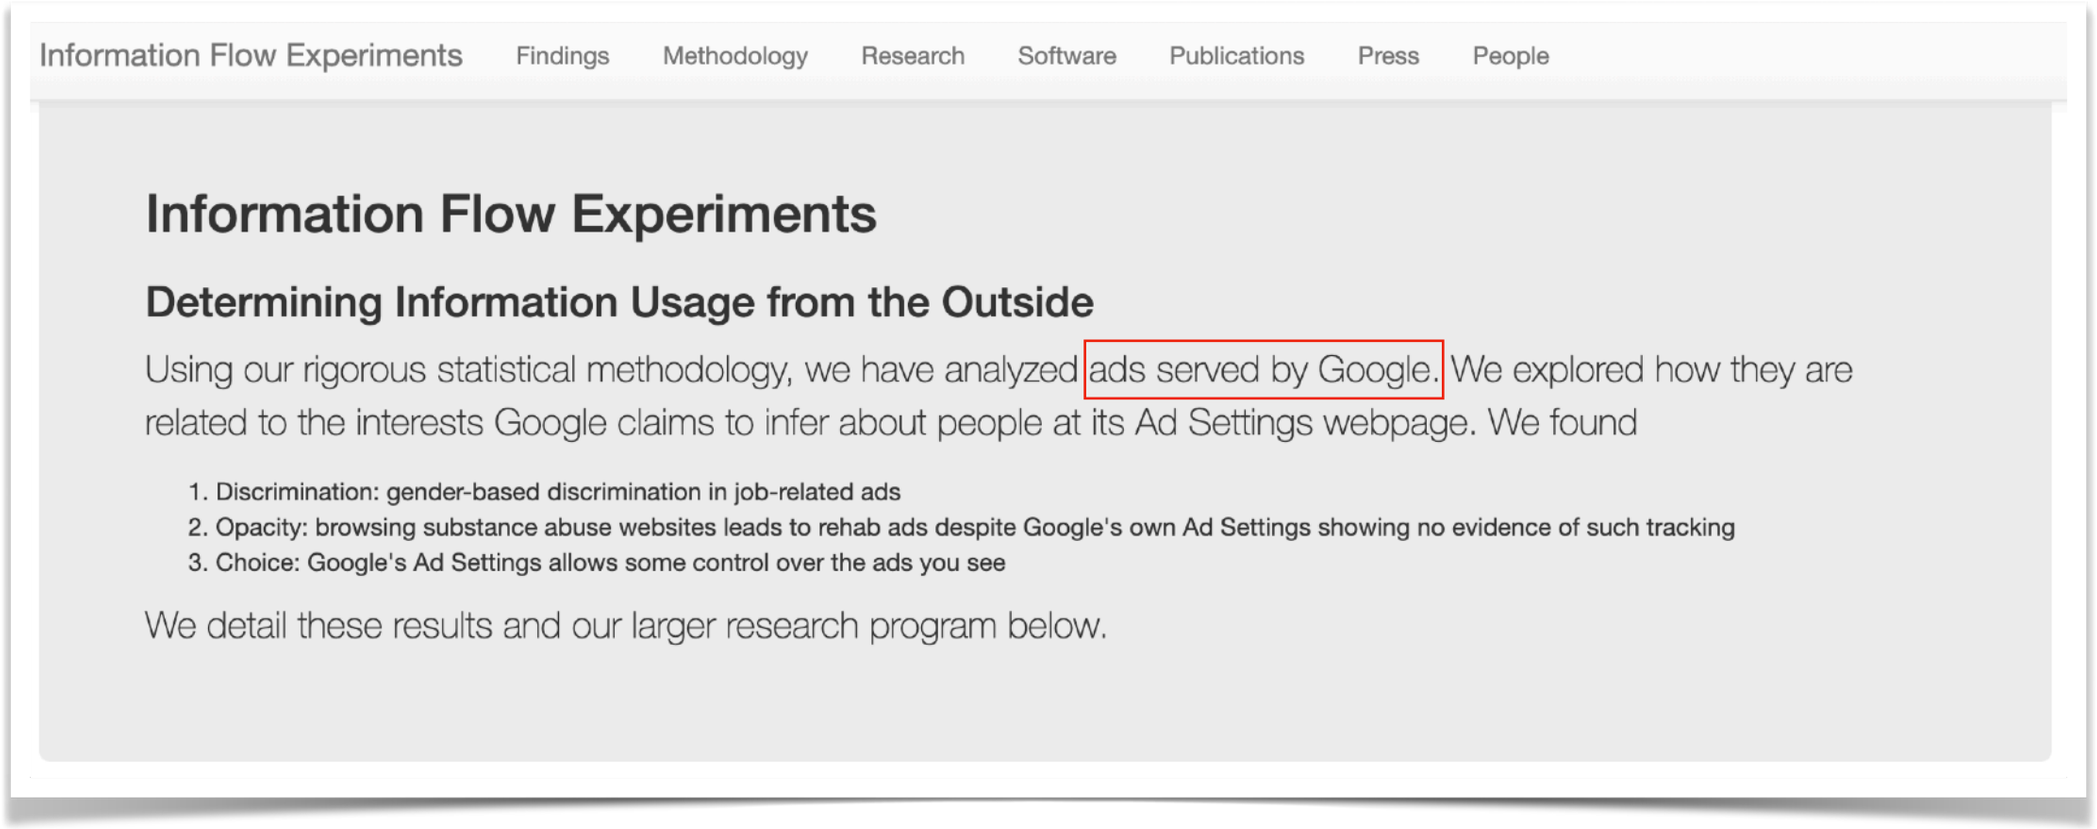

https://www.cs.cmu.edu/~mtschant/ife/

Example

Over hundreds of browsers, we randomly edited the profile to be either “female” or “male” and visited job-related websites. We found that the “male” instances were much more likely to receive ads promoting high paying jobs than the “female” instances.

Top ads for identifying the female group

Ad Title	Ad URL	Times shown to Females	Times shown to Males
Jobs (Hiring Now)	www.jobsinyourarea.co	45	8
4Runner Parts Service	www.westernpatoyotaservice.com	36	5
Criminal Justice Program	www3.mc3.edu/Criminal+Justice	29	1
Goodwill - Hiring	goodwill.careerboutique.com	121	39
UMUC Cyber Training	www.umuc.edu/cybersecuritytraining	38	30

Top ads for identifying the male group

Ad Title	Ad URL	Times shown to Females	Times shown to Males
$200k+ Jobs - Execs Only	careerchange.com	311	1816
Find Next $200k+ Job	careerchange.com	7	36
Become a Youth Counselor	www.youthcounseling.degreeleap.com	0	310
CDL-A OTR Trucking Jobs	www.tadrivers.com/OTRJobs	0	8
Free Resume Templates	resume-templates.resume-now.com	8	10

The human factor

Human Biases

Our brains are evolved to help us survive. That means they take a lot of shortcuts to help us get through the day. These shortcuts, or **heuristics, are vital. But they come at a cost**.

Our world is much more complex than the world our brains developed these heuristics. **Unconscious brains can be unreliable in this environment**.

Our unconscious can helps us in some situations, but it is not always the right tool. We must be sure that it will not hurt others.

The halo effect

People who looks healthy or attractive are also competent and good.

Reading:

Physiognomy's New Clothes, by Blaise Agüera y Arcas, Margaret Mitchell and Alexander Todorov.

Unconscious Human Biases

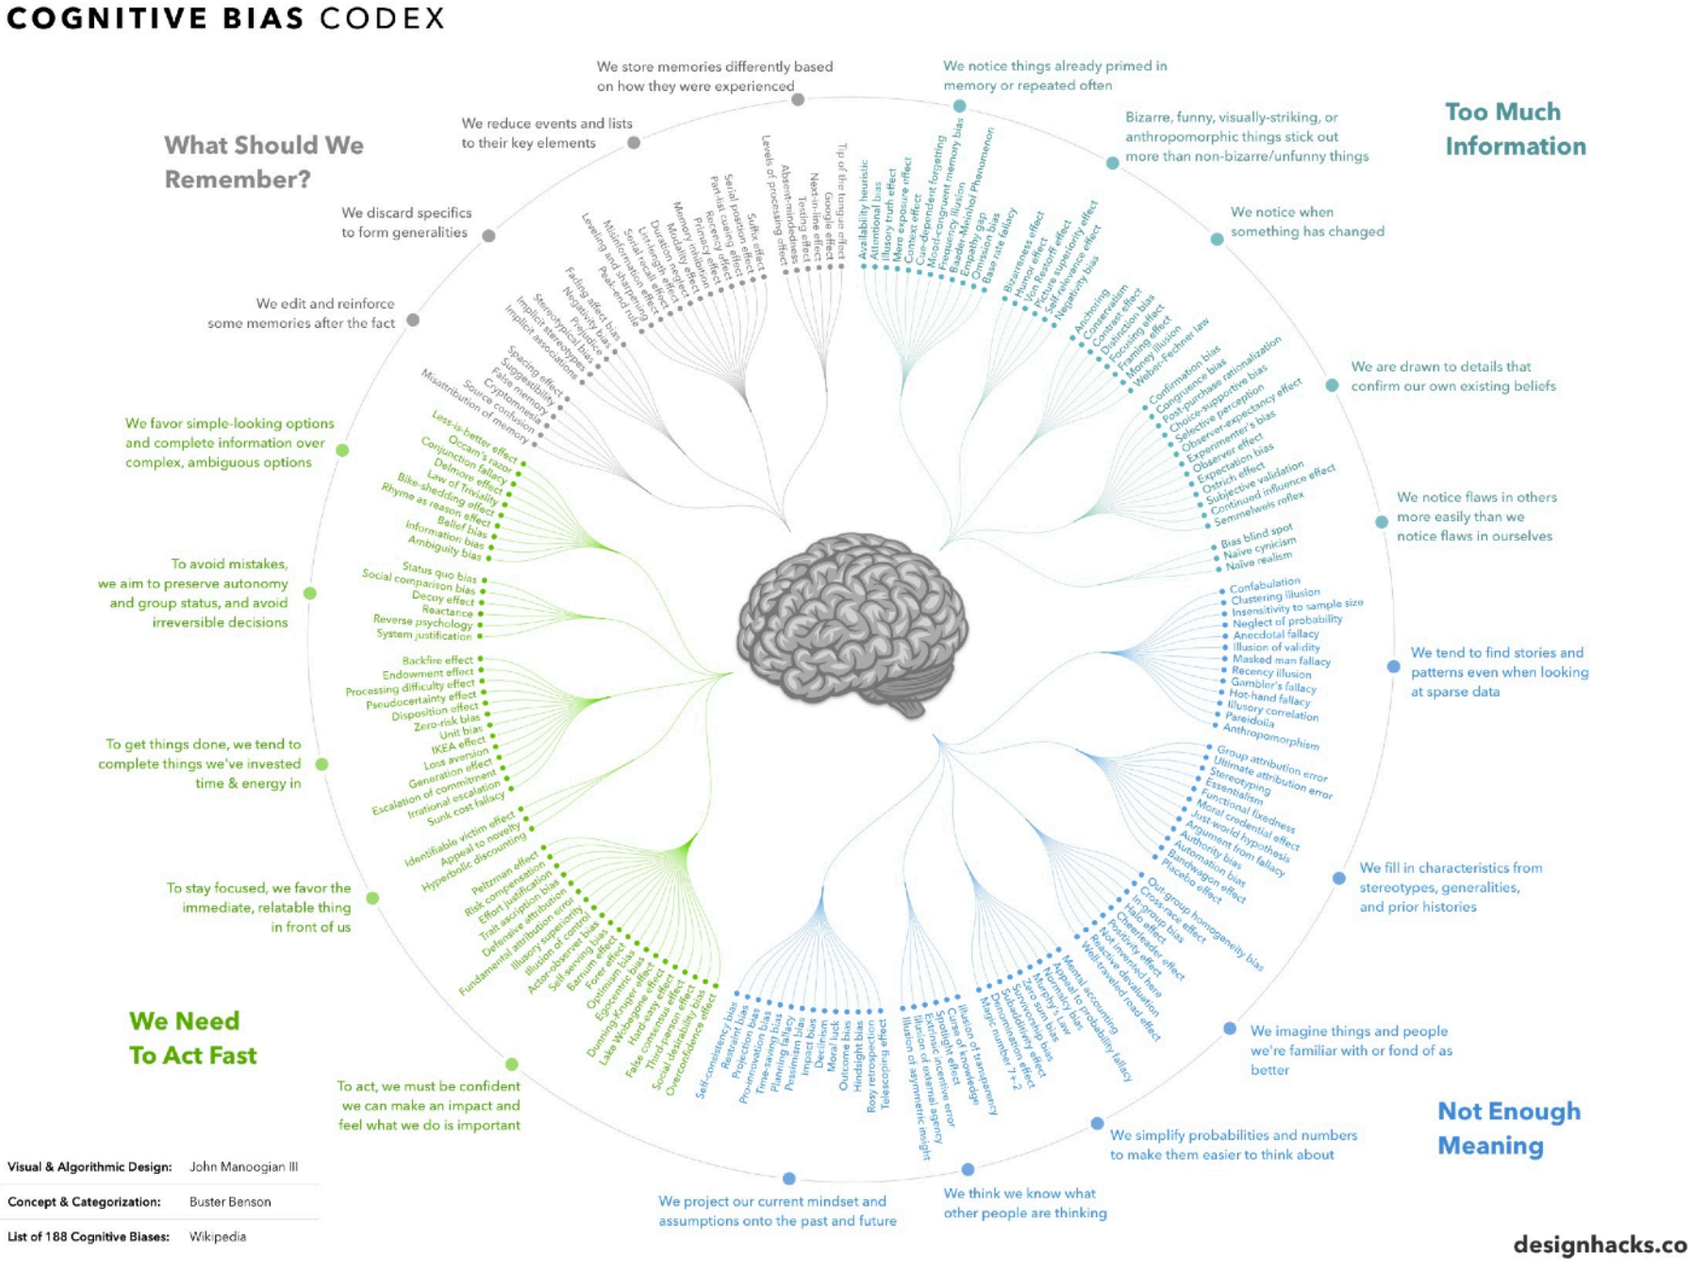

https://www.visualcapitalist.com/wp-content/uploads/2017/09/cognitive-bias-infographic.html

Human decision making

We know that human decision-making is affected by:

- Unconscious thoughts, biases, etc. MIND
- Unthinking custom and practice, or unconsciously absorving beliefs of our friends, family, society, etc. PERSONAL HISTORY
- Personal ethical decision making profile. F.e. you prioritize relationships in your decicion-making. DEFAULT SETTING
- Reflective practice, to consider context and the people who will be affected by your decisions.

The role of ethics is to have a toolkit to do reflective practice, and to able of making and justifiying our decisions

Automated Discrimination

Demographic disparities

Amazon uses a data-driven system to determine the neighborhood in which to offer free same-day delivery. A 2016 study found ***disparities in the demographic makeup of these neighborhoods****: in many U.S. cities, white residents were more than twice as likely as black residents to live in one of the qualifying neighborhoods.*

> This doesn't imply that the designer of the system intended for such inequalities to arise.
>
> Looking beyond intent, it's important to understand when observed disparities can be considered to be discrimination.

To understand why the racial disparities in Amazon's system might be harmful, we must keep in mind the history of racial prejudice in the United States, its relationship to geographic segregation and disparities, and the perpetuation of those inequalities over time.

Demographic disparities

"ENGINEER" POINT OF VIEW: THAT'S NOT MY BUSINESS!

Amazon argued that its system was justified because it was designed based on efficiency and cost considerations and that race wasn't an explicit factor.

THIS COULD BE ILLEGAL!

Nonetheless, it has the effect of **providing different opportunities to consumers at racially disparate rates**.

THIS IS UNETHICAL

The concern is that this might contribute to the perpetuation of long-lasting cycles of inequality.

THIS IS ALSO UNETHICAL, BUT AT A DIFFERENT LEVEL

If, instead, the system had been found to be partial to ZIP codes ending in an odd digit, it would not have triggered a similar outcry.

Demographic disparities

The term ~~**bias**~~ is often used to refer to **demographic disparities** in algorithmic systems that are objectionable for societal reasons.

We should avoid using the word bias since it means different things to different people.

> Suppose that Amazon's estimates of delivery dates/times were consistently too early by a few hours. This would be a case of *statistical bias*.

> Suppose that Amazon's estimates of delivery dates/times are based on the prediction of a recurrent neural network. This would be a case of *recency bias*, a kind of *inductive bias* in machine learning.

Demographic disparities

Demographic disparities propagates through the machine learning pipeline. Studying the stages of machine learning is crucial if we want to **intervene** to **minimize disparities**.

1. The first stage is **measurement**, which is the process by which the state of the world is reduced to a set of rows, columns, and values in a dataset.

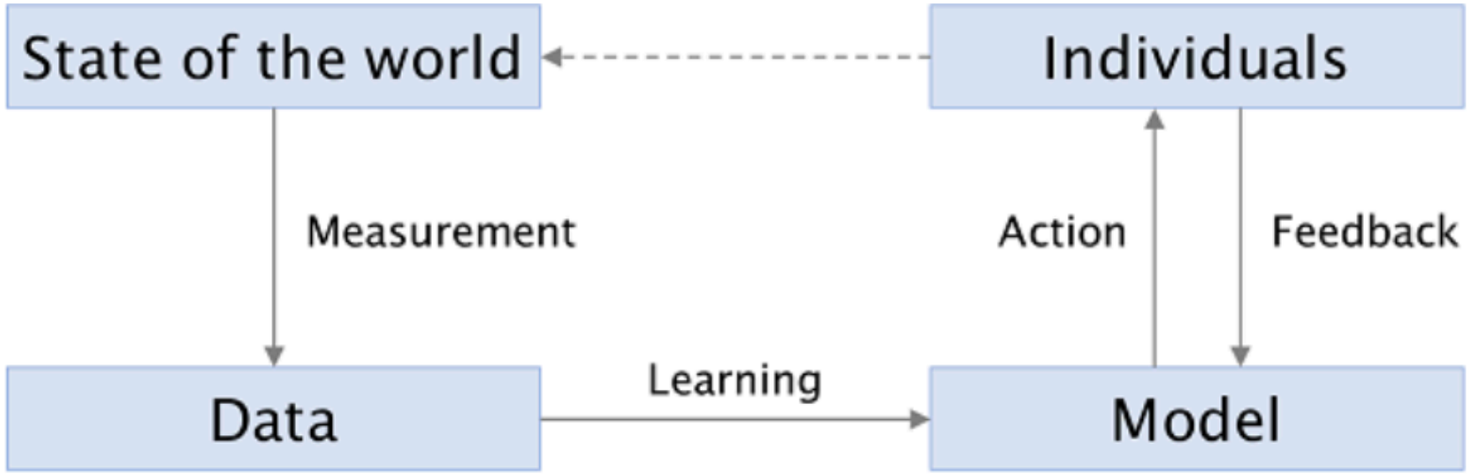

2. The '**learning**' in machine learning refers to the next stage, which is to turn that data into a model. Model summarizes the patterns in the training data; it makes generalizations.

3. The next stage is the **action** we take based on the model's predictions, which are applications of the model to new, unseen inputs.

4. Some machine learning systems record feedback from users (how users react to actions) and use them to refine the model. Feedback can also occur unintentionally, or even adversarially; these are more problematic

https://fairmlbook.org/

Demographic disparities

We're mainly concerned with applications of machine learning that involve data about **people**.

In these applications, the available training data will likely encode the demographic disparities that exist in our society.

What about applications that aren't about people?

Consider "Street Bump," a project by the city of Boston to crowdsource data on potholes.
The smartphone app automatically detects **pot holes** using data from the smartphone's sensors and sends the data to the city. Infrastructure seems like a comfortably boring application of data-driven decision-making, far removed from the ethical quandaries we've been discussing.
But the data reflects **terms of smartphone ownership**, which are higher in wealthier parts of the city compared to lower-income areas and areas with large elderly populations.

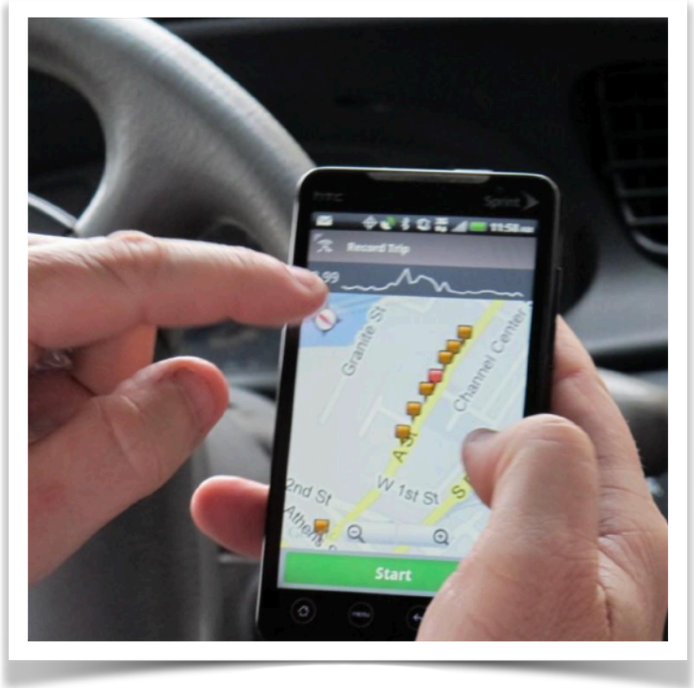

Demographic disparities

There are also troubles related to **measurement**: measuring almost any **attribute about people** is similarly subjective and challenging: *teacher effectiveness*, *economic status*, etc.

Recommended Reading:
Measurement and Fairness, by Abigail Z. Jacobs, Hanna Wallach
https://arxiv.org/abs/1912.05511

Demographic disparities

We've seen that training data reflects the disparities, distortions, and biases from the real world and the measurement process.

This leads to an obvious question: **when we learn a model from such data, are these disparities preserved, mitigated, or exacerbated?**

Some **patterns** in the training data ("smoking is associated with cancer") represent **knowledge** that we wish to mine using machine learning, while other patterns ("girls like pink and boys like blue") represent **stereotypes** or **bad habits** that we might wish to avoid learning.

But learning algorithms have no general way to distinguish between these two types of patterns, because they are the result of social norms and moral judgments.

Sensible characteristics

In many classification tasks, the features contain or implicitly encode **sensitive characteristics** of an individual.

THE PROBLEM

The choice of sensitive attributes will generally have profound consequences as it decides which groups of the population we highlight, and what conclusions we draw from our investigation.

THE BAD SOLUTION

Some have hoped that **removing or ignoring sensitive attributes** would somehow ensure the impartiality of the resulting classifier. Unfortunately, this practice is usually somewhere on the spectrum between ineffective and harmful.

THE CAUSE OF THE BAD SOLUTION

In a typical data set, we have many features that are slightly correlated with the sensitive attribute. However, if numerous such features are available, as is the case in a typical browsing history, the task of predicting gender becomes feasible at high accuracy levels.

ENGINEER POINT OF VIEW: THAT'S NOT MY BUSINESS!

But, isn't discrimination the very point of machine learning?

Yes, but it is not admissible when this discrimination/differentiation is based on unjustified causes, is practically irrelevant or is morally wrong.

WE NEED A CASE BY CASE ANALYSIS
FAIRNESS CANNOT BE AUTOMATED

Discrimination is not a general concept, it's **domain and feature specific**!

Formal non-discrimination

$X \in \mathbb{R}^d$ features of an individual (*browsing history*)

$A \in \{0,1\}$ sensitive attribute (*gender*)

$Y \in \{0,1\}$ target variable

Decision function (i.e. ML algorithm) is any random variable $D = d(X, A) \in [0,1]$ that can be transformed into a hard decision by thresholding.

Formal non-discrimination

Criterion 0 (Naive):

1. **Unawarness** : Fairness is the absence of the protected attribute in the model features.

$$d(X, A) \rightarrow d(X)$$

Removing the protected attribute doesn't guarantee that all the information concerning this attribute is removed from the data!

Formal non-discrimination

2 important criteria:

Sometimes called **demographic parity** or **statistical parity**.

1. **Independence**: D independent of A, denoted $D \perp A$. For all groups a, b and all values d:

The sensitive characteristic must be statistically independent of the score.

$$p(D = d \,|\, A = a) = p(D = d \,|\, A = b)$$

Sometimes called **equalized odds**.

2. **Separation**: D independent of A conditional on Y.

Independence

X_1	...	A = race	Y
0	...	0	1
2	...	1	1
1	...	0	0
2	...	1	0

In the case of binary decision, independence simplifies to the condition for all groups a, b:

$$p(D = 1 \,|\, A = a) = p(D = 1 \,|\, A = b)$$

Acceptance rate must be the same in all groups

Independence

Let's assume we're building an application to select promising candidates for a job. Our model will aim to learn the typical profile of those who can be hired.

In this example we get demographic parity:

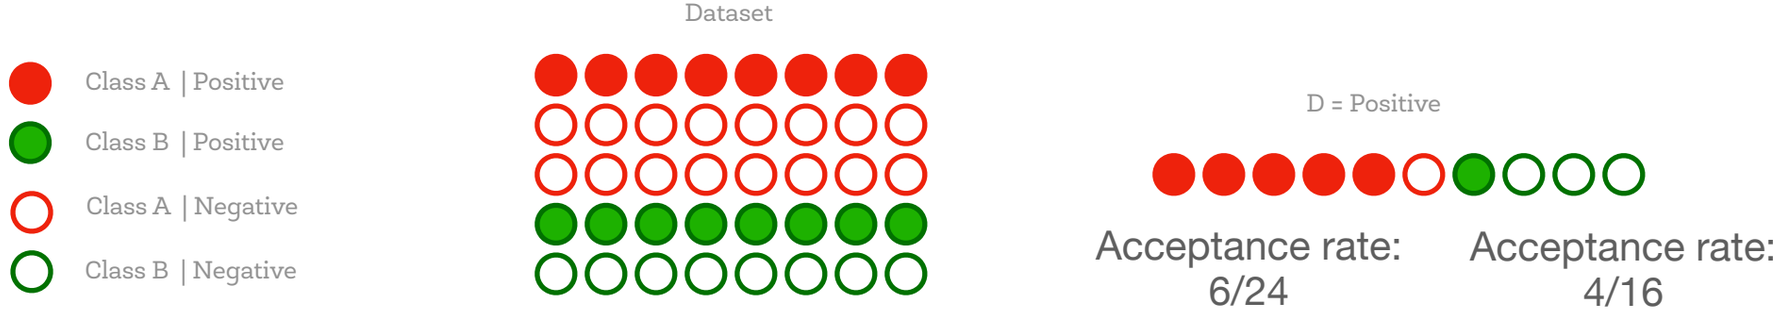

There are flaws:

- It concerns only the final outcome of the model but doesn't focus on **equality of treatment** (see next slide).
- Demographic parity can reject the optimal classifier.

Warning!

Decisions based on a classifier that satisfies independence can have **undesirable properties** (and similar arguments apply to other statistical criteria).

> Imagine a company that in group A hire diligently selected applicants at some rate $p > 0$.
>
> In group B, the company hires carelessly selected applicants at the same rate p.
>
> Eventhough the acceptance rates in both groups are identical, it is far more likely that unqualified applicants are selected in one group than in the other.
>
> As a result, it will appear in hindsight that members of group B performed worse than members of group A, thus establishing a negative track record for group B.

Independence

In general, there are three different strategies to satisfy fairness criteria:

- **Pre-processing**: Adjust the feature space to be uncorrelated with the sensitive attribute.

- At **training time**: Work the constraint into the optimization process that constructs a classifier from training data.

- **Post-processing**: Adjust a learned classifier so as to be uncorrelated with the sensitive attribute.

Achieving independence with a representation learning approach:

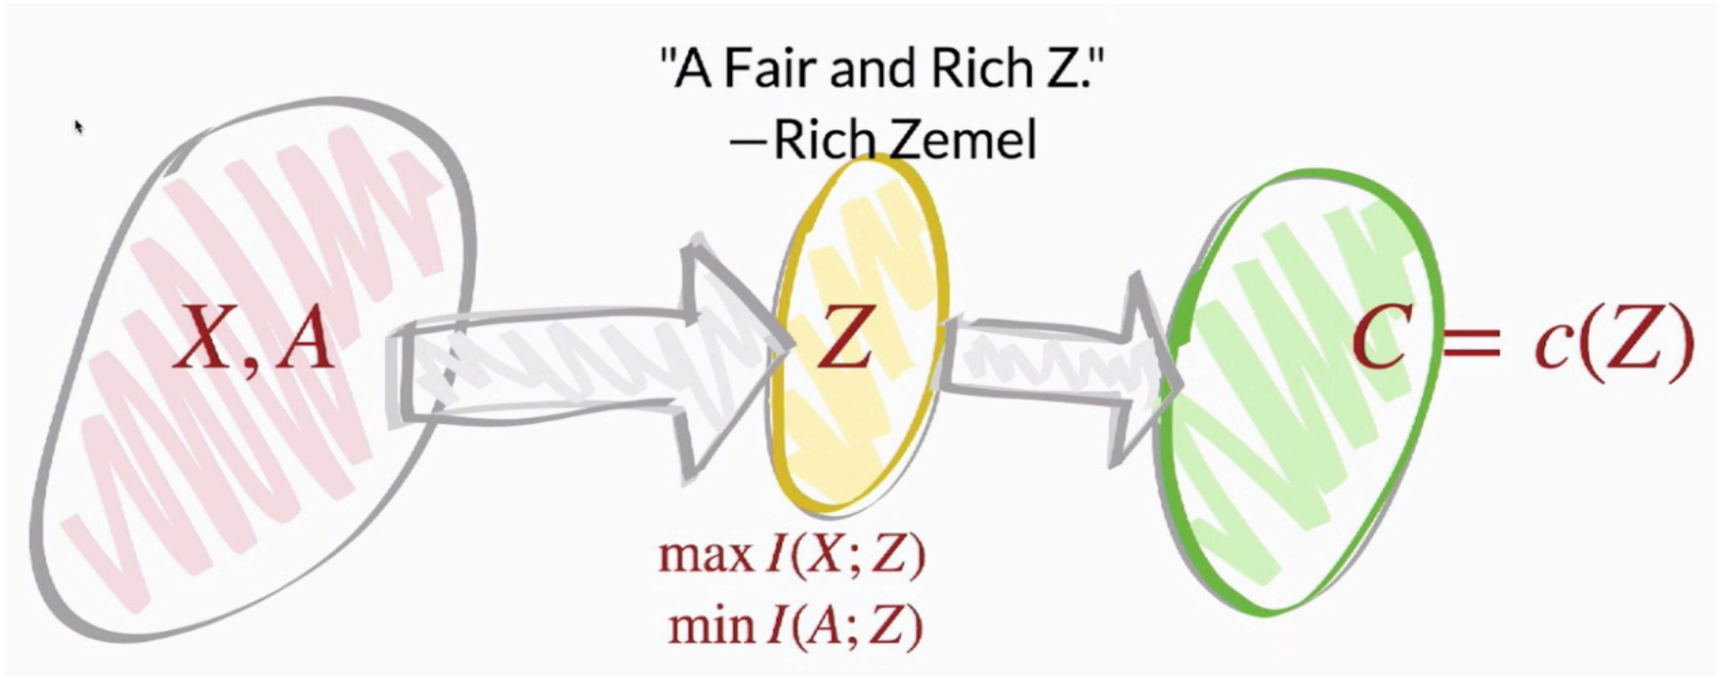

Separation

D independent of A conditional on Y.

A bank might argue that it is a matter of business necessity to therefore have different lending rates for these groups.

Separation acknowledges that in many scenarios, the **sensitive characteristic may be correlated with the target variable**. For example, one group might have a higher default rate on loans than another.

Roughly speaking, the separation criterion allows correlation between the score and the sensitive attribute to the extent that it is justified by the target variable.

Equalized odds requires both the fraction of non-defaulters that qualify for loans and the fraction of defaulters that qualify for loans to be constant across groups.

Separation

In the case where D is a binary classifier, separation is equivalent to requiring for all groups a, b the two constraints

$$p(D = 1 \mid Y = 1, A = a) = p(D = 1 \mid Y = 1, A = b)$$ True Positive Rate

$$p(D = 1 \mid Y = 0, A = a) = p(D = 1 \mid Y = 0, A = b)$$ False Positive Rate

What separation therefore requires is that all groups experience the same **false negative rate** and the **same false positive rate**.

FNR = 1 - TPR

Separation

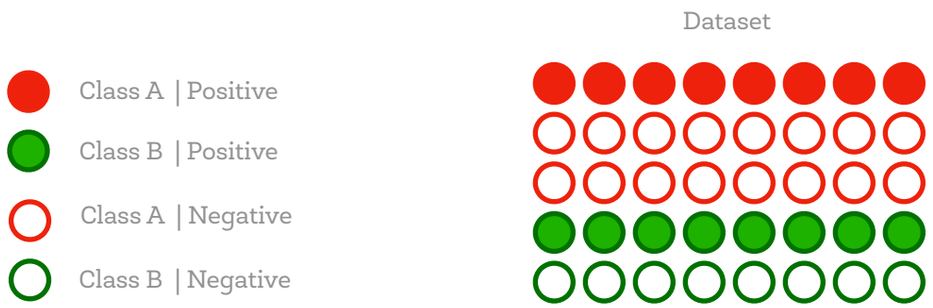

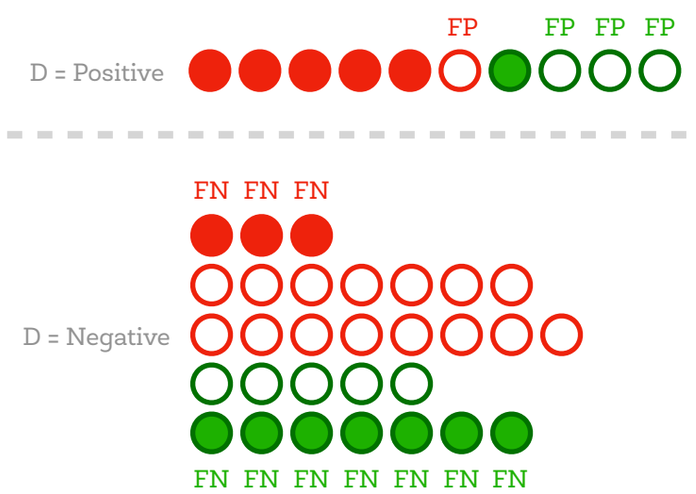

$$FPR_A = \frac{1}{16}$$

$$FPR_B = \frac{3}{8}$$

$$FNR_A = \frac{3}{8}$$

$$FNR_A = \frac{7}{8}$$

Separation

As was the case with independence, we can achieve separation by post-processing a given score function without the need for retraining.

In many applications (e.g. hiring), people care more about the true positive rate than true negative rate so many works focus on the following relaxed version (which is called **equality of opportunity**):

$$p(D = 1 \mid Y = 1, A = 0) = p(D = 1 \mid Y = 1, A = 1)$$

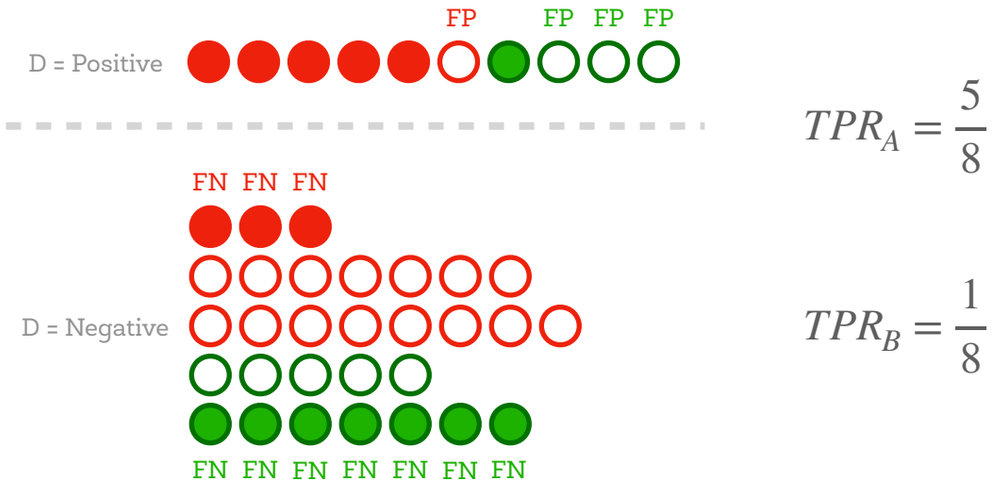

Separation flaws

It may not help closing the gap between two groups.

For example, imagine group A has 100 applicants and 58 of them are qualified while group B also have 100 applicants but only 2 of them are qualified.

If the company decides to accept 30 applicants and satisfies equality of opportunities, 29 offers will be conferred to group A while only 1 offer will be conferred to group B. If the job is a well-paid job, group A tends to have a better living condition and affords better education for their kids, and thus enable them to be qualified for such well-paid jobs when they grow up. The gap between group A and group B will tend to be enlarged over time.

Relationships between criteria

The criteria we reviewed constrain the joint distribution in non-trivial ways. We should therefore suspect that imposing any two of them simultaneously over-constrains the space to the point where only degenerate solutions remain.

It can be shown that if we assume that Y is binary, A is not independent of Y, and R is not independent of Y, then, independence and separation cannot both hold.

It is impossible to satisfy all definitions of group fairness, meaning that the data scientists need to choose one to refer to when starting a fairness analysis.

Fairness for decisions

For binary decision procedures, we can summarize a procedure with the confusion matrix, which illustrates match and mismatch between decision d and true status Y.

	Positive Status $Y=1$	**Negative Status** $Y=0$	**Prevalence ("base rate")** $P[Y=1]$	
Positive Decision $d=1$	**True Positive (TP)**	**False Positive (FP)**	**Positive Predictive Value (PPV), aka precision** $P[Y=1\|d=1]$	**False Discovery Rate (FDR)** $P[Y=0\|d=1]$
Negative Decision $d=0$	**False Negative (FN)**	**True Negative (TN)**	**False Omission Rate (FOR)** $P[Y=1\|d=0]$	**Negative Predictive Value (NPV)** $P[Y=0\|d=0]$
Positive Decision Rate $P[d=1]$	**True Positive Rate (TPR), aka recall, aka sensitivity** $P[d=1\|Y=1]$	**False Positive Rate (FPR)** $P[d=1\|Y=0]$	**Accuracy** $P[d=Y]$	
	False Negative Rate (FNR) $P[d=0\|Y=1]$	**True Negative Rate (TNR), aka specificity** $P[d=0\|Y=0]$		

Confusion Matrix

For any box in the confusion matrix involving the decision d, we can define fairness as equality across groups.

For example, **Equal False Omission Rates**:

$p(Y=1\,|\,d=0,A=a)=p(Y=1\,|\,d=0,A=a')$

Fairness for decisions

For binary decision procedures, we can summarize a procedure with the confusion matrix, which illustrates match and mismatch between decision d and true status Y.

	Positive Status $Y = 1$	**Negative Status** $Y = 0$	**Prevalence ("base rate")** $P[Y = 1]$	
Positive Decision $d = 1$	**True Positive (TP)**	**False Positive (FP)**	**Positive Predictive Value (PPV), aka precision** $P[Y = 1 \mid d = 1]$	**False Discovery Rate (FDR)** $P[Y = 0 \mid d = 1]$
Negative Decision $d = 0$	**False Negative (FN)**	**True Negative (TN)**	**False Omission Rate (FOR)** $P[Y = 1 \mid d = 0]$	**Negative Predictive Value (NPV)** $P[Y = 0 \mid d = 0]$
Positive Decision Rate $P[d = 1]$	**True Positive Rate (TPR), aka recall, aka sensitivity** $P[d = 1 \mid Y = 1]$	**False Positive Rate (FPR)** $P[d = 1 \mid Y = 0]$	**Accuracy** $P[d = Y]$	
	False Negative Rate (FNR) $P[d = 0 \mid Y = 1]$	**True Negative Rate (TNR), aka specificity** $P[d = 0 \mid Y = 0]$		

Confusion Matrix

Demographic Parity:
$p(D = 1 \mid A = a) = p(D = 1 \mid A = a')$

Fairness for scores

For score outputs, we can consider the following initial definitions of fairness based on equal metrics across groups:

- **Balance for the Positive Class**: the average score assigned to positive members, $\mathbb{E}(D \mid Y = 1)$, should be the same across groups.

- **Balance for the Negative Class**: the average score assigned to negative members, $\mathbb{E}(D \mid Y = 0)$, should be the same across groups.

- **Calibration**: the fraction of those marked with a given score who are actually positive, $\mathbb{E}(Y \mid D = d)$, should be the same across groups.

- **AUC (Area Under Curve) Parity**: the area under the receiver operating characteristic (ROC) curve should be the same across groups. The AUC can be interpreted as the probability that a randomly chosen positive individual $Y = 1$ is scored higher than a randomly chosen negative individual.

Calibration

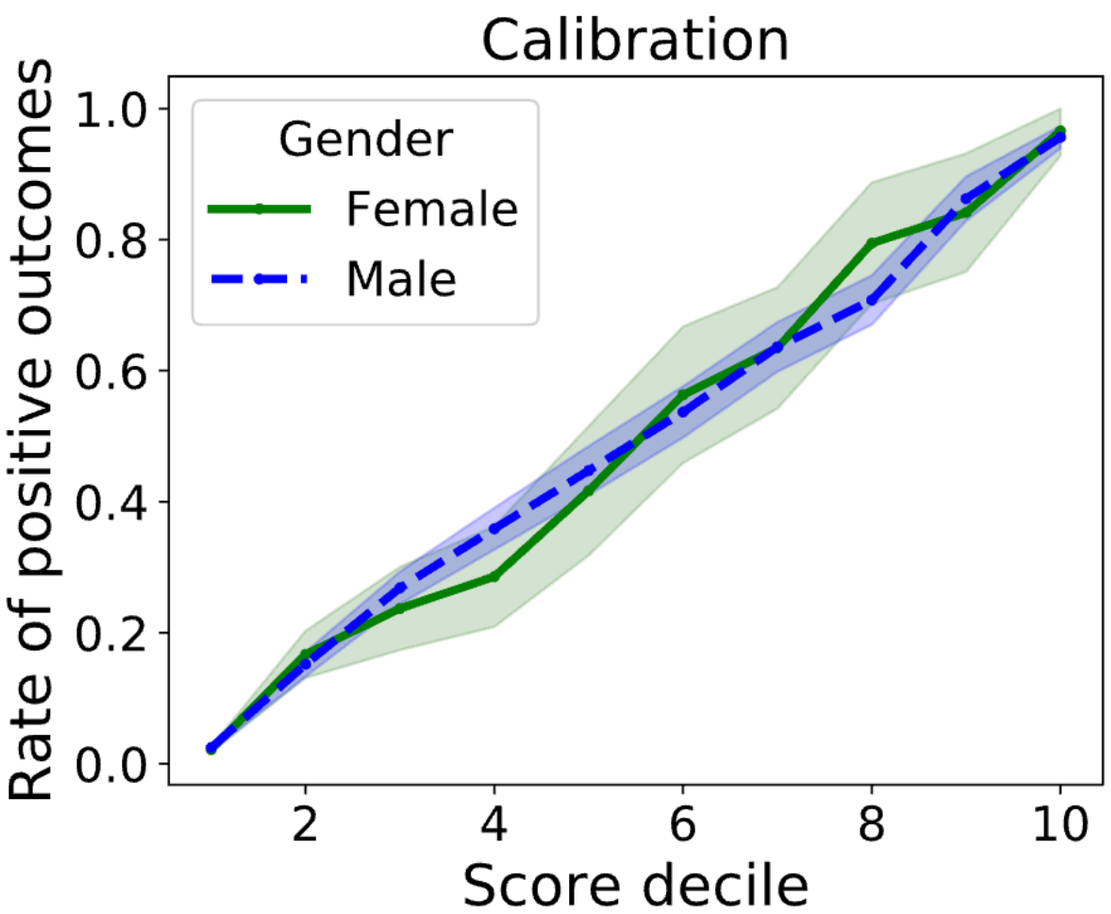

Calibration by gender on UCI adult data. A straight diagonal line would correspond to perfect calibration.

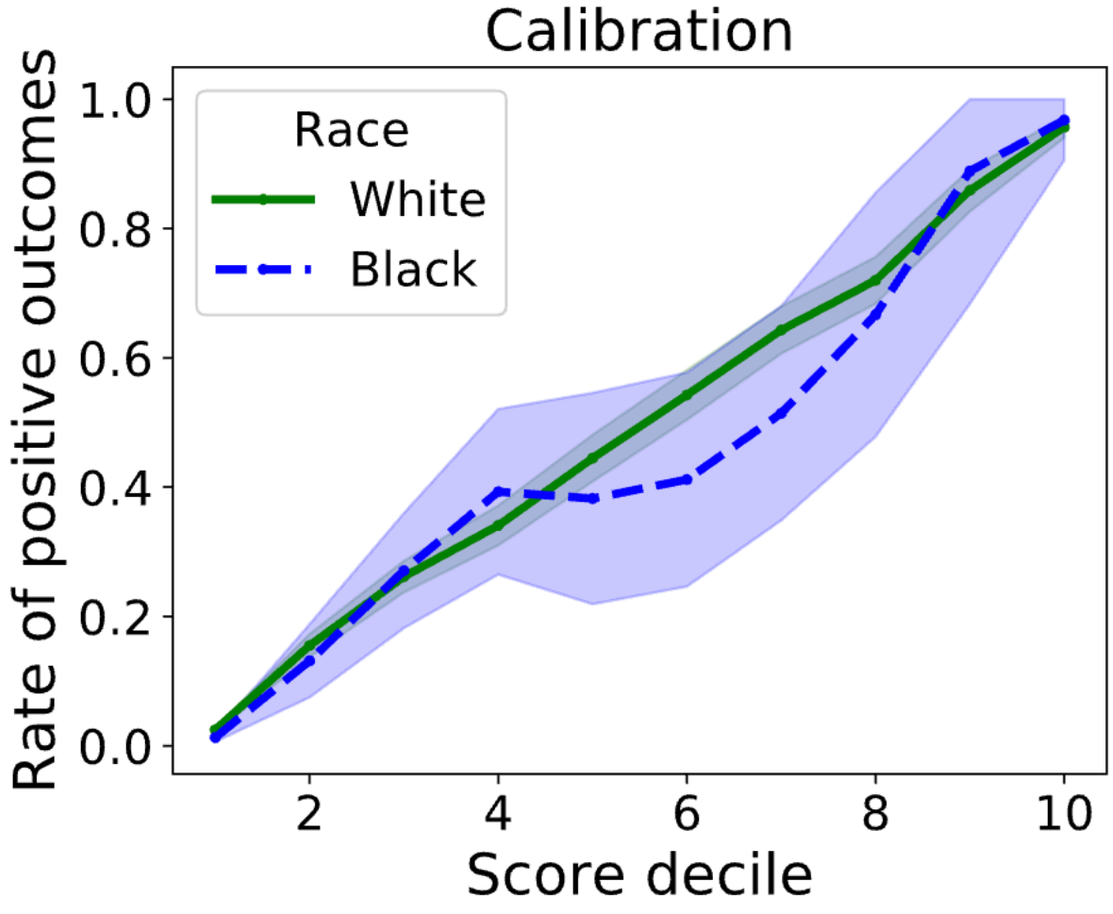

Calibration by race on UCI adult data.

Conclusions

Fairness is a not a property of a model. It is a never-ending process to align the behavior of the model and its consequences to our changing principles and values.

Fairness is a highly contextual problem. **Which fairness metric should we use to measure disparity?**

Conclusions

To help answer this question, we can define two types of fairness metrics.

'**Bias preserving**' fairness metrics seek to reproduce historic performance in the outputs of the target model with equivalent error rates for each group as reflected in the training data (or status quo).

In contrast, '**bias transforming**' metrics do not blindly accept social bias as a given or neutral starting point that should be preserved, but instead require people to make an explicit decision as to which biases the system should exhibit.

Bias Preservation in Machine Learning:
The Legality of Fairness Metrics Under EU Non-Discrimination Law

Sandra Wachter[1], Brent Mittelstadt[2] and Chris Russell[3]

Conclusions

We say that any fairness metric is bias preserving if it is always satisfied by a perfect classifier that exactly predicts its target labels with zero error, replicating bias present in the data.

Fairness metrics that are not necessarily satisfied by a perfect classifier, we refer to as bias transforming.

For example, equalized odds is bias preserving metrics: it is a form of conditional independence or conditional demographic parity, conditioned on historic data Y, and reflecting its biases exactly.

Bias Preservation in Machine Learning:
The Legality of Fairness Metrics Under EU Non-Discrimination Law

Sandra Wachter[1], Brent Mittelstadt[2] and Chris Russell[3]

Conclusions

Fairness metric	Bias preserving?
1. Group fairness, Statistical (demographic) parity	✗
2. Conditional statistical (demographic) parity, Conditional independence	✗
3. Predictive parity, outcome test	✓
4. False positive error rate balance	✓
5. False negative error rate balance, Equal opportunity	✓
6. Equalized odds	✓
7. Conditional use accuracy equality	✓
8. Overall accuracy equality	✓
9. Treatment equality	✓
10. Test-fairness or calibration	✓
11. Well-calibration	✓
12. Balance for positive class	✓
13. Balance for negative class	✓
14. Causal discrimination (direct discrimination)	*
15. Fairness through unawareness	*
16. Fairness through awareness	✗
17. Counterfactual fairness	✗
18. No unresolved discrimination	✗
19. No proxy discrimination	✗
20. Path based causal reasoning	✗

Table 1 – Bias preserving fairness metrics

Bias Preservation in Machine Learning:
The Legality of Fairness Metrics Under EU Non-Discrimination Law

Sandra Wachter[1], Brent Mittelstadt[2] and Chris Russell[3]

Conclusions

The key difference between **bias transforming** and **bias preserving** metrics is that most bias transforming metrics are satisfied by matching **decision rates between groups**, while bias preserving metrics typically require matching **error rates between groups**.

Conclusions

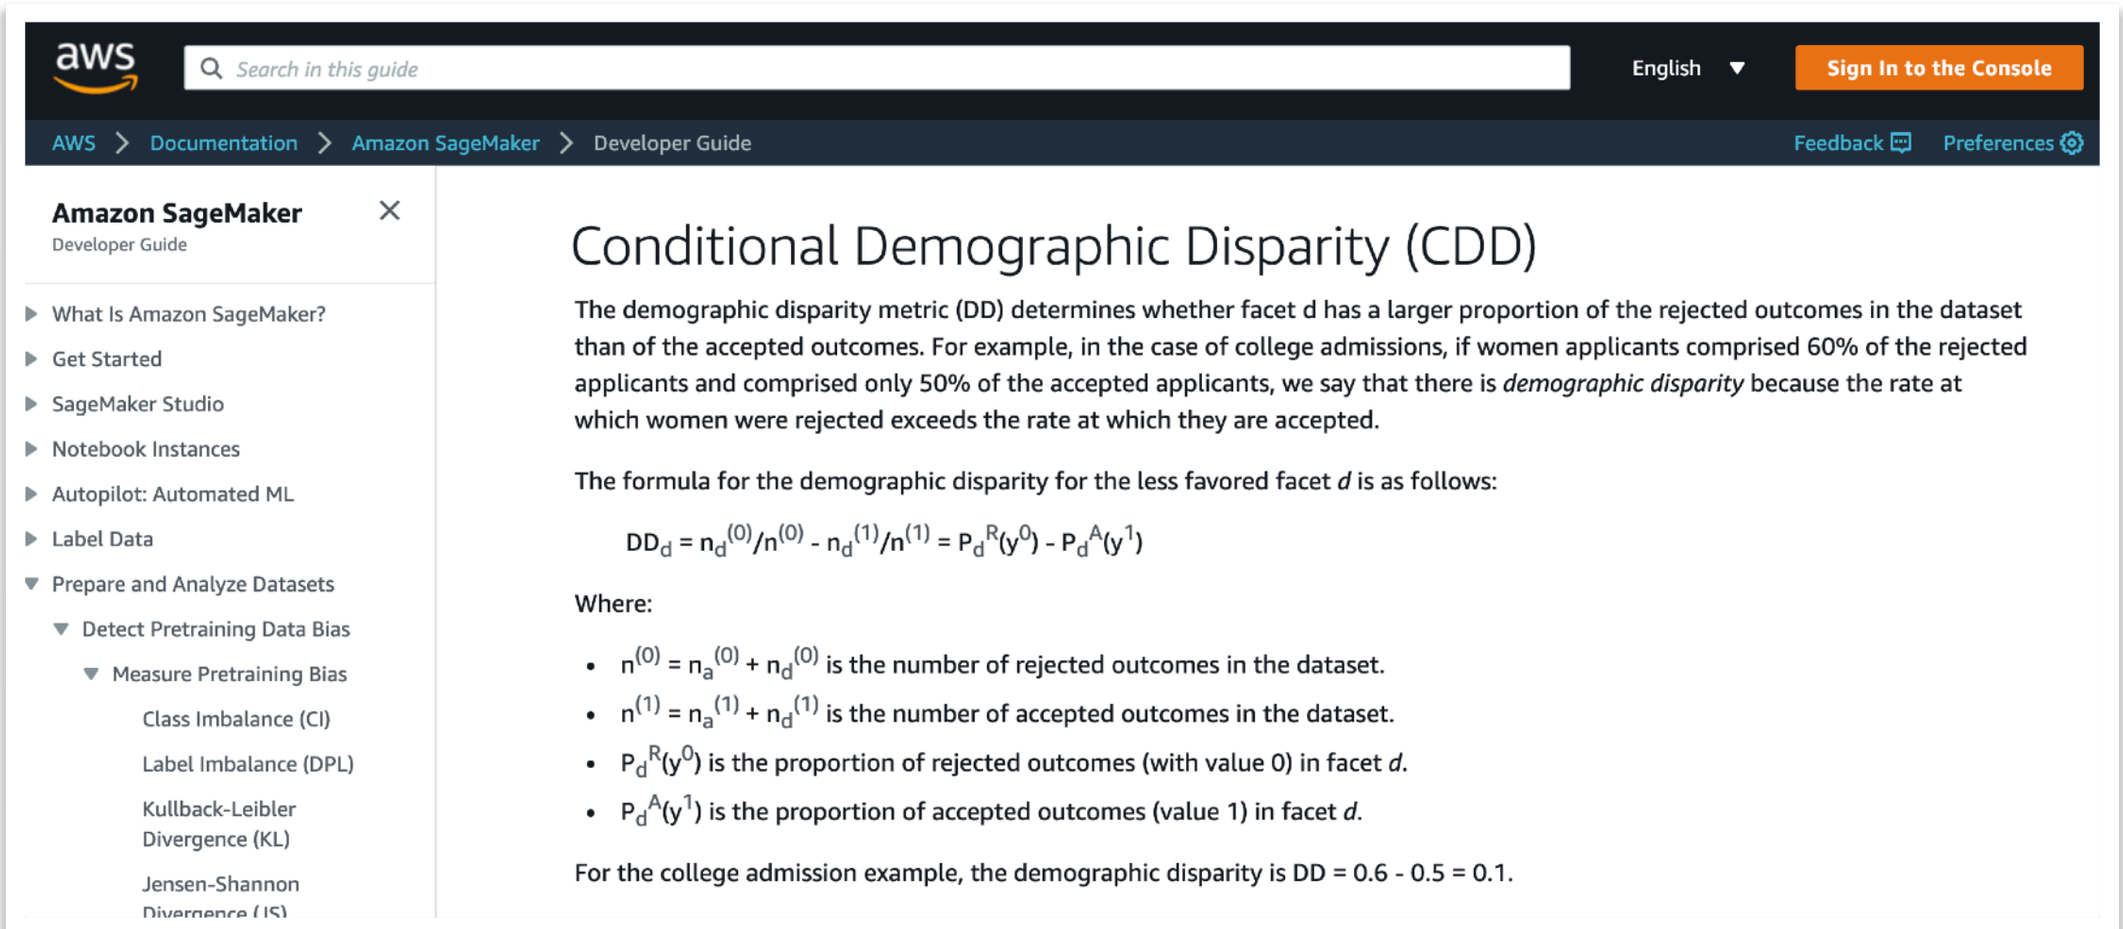

Conditional demographic parity (conditioning on salary) is satisfied if x% of both black and white people earning over a threshold receive positive decisions and also some y% of both black and white people earning under the threshold receive positive decisions.

Measuring recidivism risk

The criminal justice system needs to evaluate a diverse set of **risks**:

- The risk of committing a new crime (recidivism),
- The risk of committing a new violent crime (violent recidivism),
- The risk of committing an act of violence against another inmate or penitentiary personnel in jail (intra-penitentiary violence),
- The risk of committing an administrative violation such as breaking the conditions of a permit.

Measuring recidivism risk

Structured risk assessment corresponds to a family of methodologies for evaluating these risks using a systematic process, typically in which a number of different items are evaluated, and a final decision is taken based on the results of those items.

We can train a ML system to make automatic decisions based on the scores in each item, bu most often, a professional makes a decision based on his/her own evaluation of a defendant and the result of a series of items.

Measuring recidivism risk

Between arrest and trial, US criminal law leaves a lot of latitude to interpretation of whether accused individuals should be incarcerated, in what kind of facility, and for how long.

There has been an increasing interest in either automating these decisions or including automated insights in decision making, in the form of risk scoring for “recidivism” or “re-offense” (one can object to both terms as creating an impression of certainty that there was a first offense).

Measuring recidivism risk

COMPAS (Correctional Offender Management Profiling for Alternative Sanctions) is an automatic tool that outputs numerical scores, which are labeled, for example, “risk of recidivism”, “risk of violent recidivism”, or “risk of failure to appear”.

These scores are then used in an unspecified way to make decisions of jail, bail, home arrest, release, etc.

Measuring recidivism risk

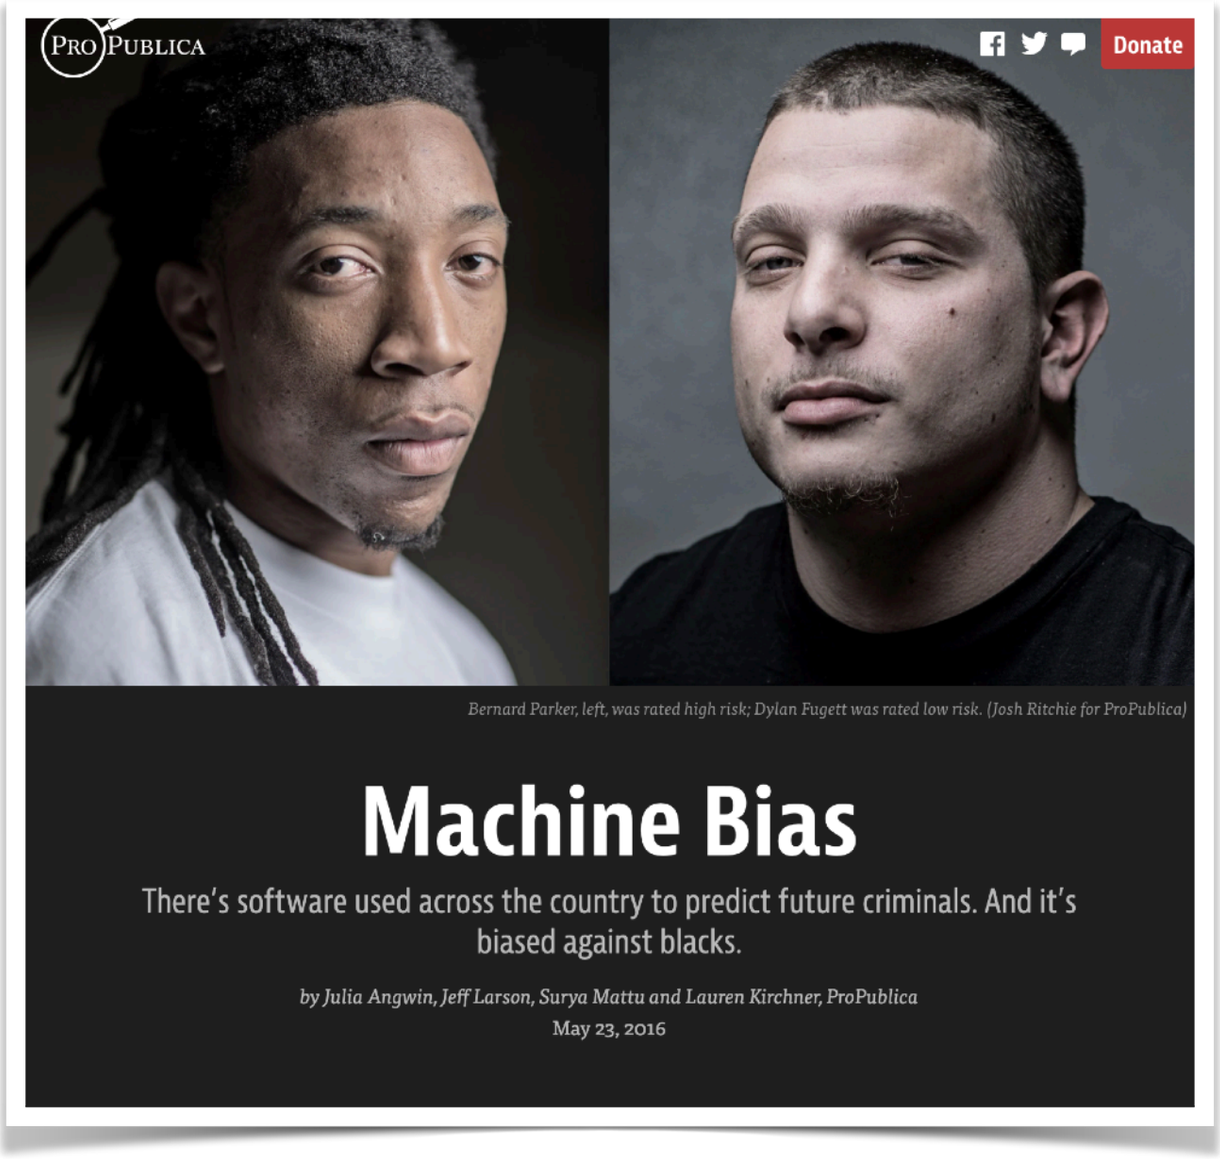

In May 2016, ProPublica, an investigative journalism organizations, published a piece called "Machine Bias", in which a predictive tool for predicting recidivism used in the US, **COMPAS**, was found **to be biased** against African-Americans.

Measuring recidivism risk

ProPublica Local Initiatives Data Store

Donate

PROPUBLICA Graphics & Data Newsletters About

Get the Big Story Join

Racial Justice Technology Criminal Justice Regulation More... Series Video Impact Search

MACHINE BIAS

Bias in Criminal Risk Scores Is Mathematically Inevitable, Researchers Say

ProPublica's analysis of bias against black defendants in criminal risk scores has prompted research showing that the disparity can be addressed — if the algorithms focus on the fairness of outcomes.

by Julia Angwin and Jeff Larson, Dec. 30, 2016, 4:44 p.m. EST

Ad

Resolució PACs PECs. Aprueba sin Esfuerzo. Consigue soluciones rápidas y de calidad.

Pacsolver Pedir precios

FOLLOW PROPUBLICA

Twitter Facebook

YouTube RSS

STAY INFORMED

Measuring recidivism risk: the debate

In 2016, ProPublica published a highly influential analysis based on data obtained through public records requests. Two of their findings can be phrased in our language as follows:

- COMPAS does not satisfy **equal false negative rates**, in fact, white defendants who did get rearrested were nearly twice as likely to be misclassified as low risk.

- COMPAS does not satisfy **equal false positive rates**, in fact, black defendants who did not get rearrested were nearly twice as likely to be misclassified as higher risk.

In their response, Equivant/Northpointe, the developers of COMPAS, cited two articles finding that:

- COMPAS satisfies calibration: scores mean the same thing regardless of the defendant's race. For example, among defendants with a score of 7, 60 percent of white defendants were rearrested and 61 percent of black defendants were rearrested.

- COMPAS satisfies equal positive predictive values: among those labeled higher risk, the proportion of defendants who got rearrested is approximately the same regardless of race.

Measuring recidivism risk: analysis

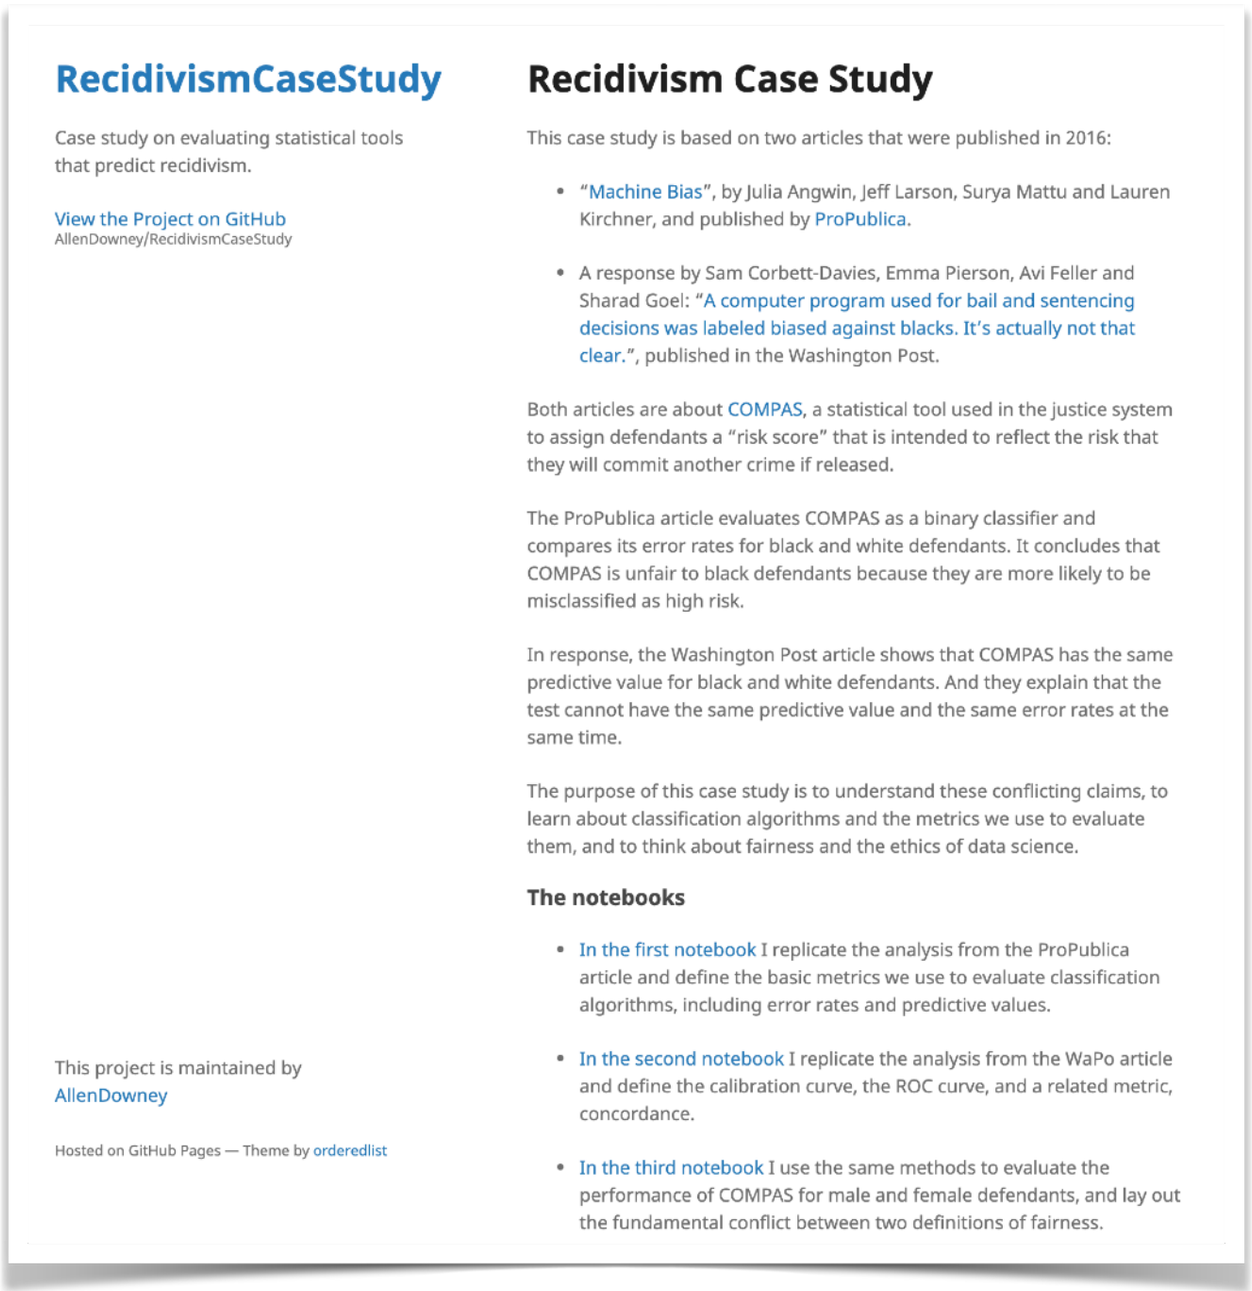

RecidivismCaseStudy

Case study on evaluating statistical tools that predict recidivism.

View the Project on GitHub
AllenDowney/RecidivismCaseStudy

This project is maintained by
AllenDowney

Hosted on GitHub Pages — Theme by orderedlist

Recidivism Case Study

This case study is based on two articles that were published in 2016:

- "Machine Bias", by Julia Angwin, Jeff Larson, Surya Mattu and Lauren Kirchner, and published by ProPublica.
- A response by Sam Corbett-Davies, Emma Pierson, Avi Feller and Sharad Goel: "A computer program used for bail and sentencing decisions was labeled biased against blacks. It's actually not that clear.", published in the Washington Post.

Both articles are about COMPAS, a statistical tool used in the justice system to assign defendants a "risk score" that is intended to reflect the risk that they will commit another crime if released.

The ProPublica article evaluates COMPAS as a binary classifier and compares its error rates for black and white defendants. It concludes that COMPAS is unfair to black defendants because they are more likely to be misclassified as high risk.

In response, the Washington Post article shows that COMPAS has the same predictive value for black and white defendants. And they explain that the test cannot have the same predictive value and the same error rates at the same time.

The purpose of this case study is to understand these conflicting claims, to learn about classification algorithms and the metrics we use to evaluate them, and to think about fairness and the ethics of data science.

The notebooks

- In the first notebook I replicate the analysis from the ProPublica article and define the basic metrics we use to evaluate classification algorithms, including error rates and predictive values.
- In the second notebook I replicate the analysis from the WaPo article and define the calibration curve, the ROC curve, and a related metric, concordance.
- In the third notebook I use the same methods to evaluate the performance of COMPAS for male and female defendants, and lay out the fundamental conflict between two definitions of fairness.

https://allendowney.github.io/RecidivismCaseStudy/

Assignment: Recividism Analysis

The dataset corresponds to a set of 4754 juvenile offenders in Catalonia who were evaluated using **SAVRY**, a structured risk assessment tool. The data on recidivism indicates if the same people committed a new offence in 2013-2015.

Objectives:

- Data Exploration
- Building a predictor
- Measuring Fairness